

KODA: Kinematics–Observation Data Alignment for Humanoid Learning from Human Egocentric Demonstration

Zibo Zhou^{1,2,3}, Yueru Chen^{2,3}, Junsong Wu^{2,3}, Weiji Xie^{2,3}, Qingyao Xu³, Sheng Yin³, Jiyuan Shi^{1,2}, Siheng Chen³, Chenjia Bai^{1,2,*}, Xuelong Li^{2,*}

¹Gamma Robotics, ²Institute of Artificial Intelligence (TeleAI), China Telecom, ³Shanghai Jiao Tong University

Training humanoid whole-body vision–language–action (VLA) policies requires costly robot demonstrations with aligned robot-view observations and controller-compatible whole-body state–action labels. Egocentric human demonstrations offer scalable supervision but have misaligned views and sparse motion cues. Existing approaches either pretrain policies that still require robot data for adaptation or generate controller-specific commands that limit reuse. Conversion also requires robot-view observations aligned with the generated whole-body motion. We present KODA (Kinematics–Observation Data Alignment), which converts egocentric RGB demonstrations into whole-body VLA training data using a reusable canonical trajectory to guide robot-view synthesis and generate controller-specific whole-body state–action supervision. KODA first generates a simulator-checked whole-body trajectory using a sparse tracker with sampling-based command optimization. This generated trajectory directly determines the onboard camera poses used to synthesize corresponding robot-view observations. The motion, observations, and task instruction form a canonical record. For each deployment controller, KODA optimizes commands to track the stored motion and records the corresponding action labels and proprioceptive states. When switching controllers, KODA retains this canonical record and regenerates whole-body state–action supervision for co-training with matching robot demonstrations. We evaluate KODA on four Unitree G1 loco-manipulation tasks using two whole-body controllers, SONIC and HEFT. With robot data held fixed, macro-average success on conditions covered by human demonstrations but absent from robot data improves by 45.0 and 46.25 percentage points, respectively. Under both controllers, this metric shows an upward trend across the evaluated human-data budgets, suggesting potential for further scaling.

Keywords: Egocentric Human Data, Humanoid Loco-Manipulation

Date: October 8, 2026

Project: <https://koda-humanoid.github.io/koda/>

Correspondence to: Chenjia Bai (baicj@chinatelecom.cn)



1 Introduction

Humanoid robots combine locomotion and dexterous manipulation in a versatile platform for operating in environments designed for people. Realizing this potential requires policies that coordinate body posture, mobility, and bimanual manipulation across the whole body. Whole-body vision–language–action (VLA) models aim to learn such coordination from onboard observations and task instructions (Jiang et al., 2026; Hu et al., 2026). Training requires corresponding robot observations, whole-body states, and action supervision compatible with the deployment whole-body controller (WBC) interface. Collecting these demonstrations on the target humanoid is costly, limiting coverage across tasks, objects, and environments.

Human demonstration interfaces and egocentric datasets provide RGB images, camera poses, and sparse bilateral hand-position references at scale (Chi et al., 2024; Hoque et al., 2026; Punamiya et al., 2026).




Figure 1 Overview of KODA. **Whole-Body Motion Generation** produces simulation-feasible whole-body motion from human egocentric demonstrations, and **Robot Onboard View Generation** synthesizes motion-aligned onboard views. The motion, views, and task instruction form a reusable canonical record. **Control Command Generation** uses this record to generate controller-specific states and actions. The resulting data are combined with matching real-robot demonstrations to co-train a controller-specific vision–language–action (VLA) policy for language-guided humanoid loco-manipulation.

Their success on fixed-base arms (Chi et al., 2024) and in dexterous manipulation (Zheng et al., 2026b; Liu et al., 2026b) motivates extending human-data supervision to whole-body VLAs. Approaches such as EgoVLA, WholeBodyVLA, and MotionWAM learn human-derived priors for humanoid policies but still depend on target-humanoid demonstrations for adaptation (Yang et al., 2025b; Jiang et al., 2026; Zheng et al., 2026a). Target-humanoid data collection therefore remains a bottleneck. Other work converts human demonstrations into robot-oriented supervision. EgoHumanoid uses whole-body human measurements and a shared end-effector-and-navigation action space for co-training (Shi et al., 2026). HALOMI trains policies on human demonstrations using head–hand command labels tied to a dedicated controller (Zhao et al., 2026). These methods produce supervision tied to particular action interfaces, creating data silos that limit reuse across whole-body controllers; their outputs do not directly provide the whole-body states and compatible actions needed for other WBC interfaces. Their converted episodes also require observations corresponding to the generated robot motion. Fixed-offset reprojection (Shi et al., 2026) shifts the human camera without following the generated robot trajectory, leaving this observation–motion correspondence unresolved. Moreover, retaining only controller-tuned commands limits reuse to that controller’s command space and tracking behavior. A reusable record must support adaptation to different deployment controllers.

We target a distinct data-conversion setting: augmenting an existing controller-specific humanoid whole-body dataset from egocentric RGB, camera poses, and hand references. This setting requires three capabilities: (i) generating simulator-checked whole-body trajectories from sparse human references; (ii) synthesizing onboard observations aligned with the generated robot motion; and (iii) instantiating the same converted record as compatible state–action supervision for different deployment controllers.

To address these three challenges, we present KODA (Kinematics–Observation Data Alignment), a framework




Figure 2 KODA pipeline. **Middle:** Human demonstrations become aligned whole-body motion and onboard observations, then controller-specific state-action data for VLA training (Control Command Generation). **Top:** Whole-Body Motion Generation combines hand estimation and contact-aware retargeting, sampling-based command optimization (SCO), and arm refinement. **Bottom:** Robot Onboard View Generation combines arm removal, scene reprojection, robot rendering, depth compositing, and inpainting. Input depth is observed or estimated from RGB.

that converts egocentric RGB demonstrations into whole-body VLA training data. The conversion proceeds in three stages, as shown in Fig. 1. First, KODA generates a whole-body trajectory from sparse human references using a fine-tuned tracker. Sampling-based command optimization improves reference tracking while maintaining trajectory feasibility in simulation. Second, KODA derives camera poses from the generated trajectory and synthesizes corresponding onboard observations. The aligned motion and observations, together with the task instruction, form a canonical record for the target humanoid. Third, KODA optimizes the selected deployment controller’s commands to track the fixed canonical trajectory, producing compatible state-action supervision for co-training with matching real-robot demonstrations. Changing controllers regenerates the state-action supervision while retaining the canonical motion and robot-view observations.

Our contributions are: (a) A data-conversion framework that constructs whole-body robot motion and corresponding onboard observations from egocentric human demonstrations. (b) A canonical representation and controller-specific control command generation procedure for generating compatible state-action supervision from the same converted records across deployment interfaces. (c) Real-world evidence of the training value and scaling potential of converted human demonstrations on four G1 loco-manipulation tasks. With 50 robot demonstrations per task fixed, adding 100 converted human demonstrations per task increases macro-average success by 45.0 and 46.25 percentage points under SONIC and HEFT, respectively, on object conditions covered by human training clips but absent from robot demonstrations. This metric shows an upward trend across the evaluated data budgets under both interfaces.

2 Related Work

2.1 Humanoid Loco-Manipulation

Humanoid data collection methods use direct teleoperation, human imitation, or whole-body tracking to obtain robot demonstrations. OmniH2O, HOMIE, HumanPlus, and TWIST2 provide interfaces for whole-

body teleoperation or imitation learning (He et al., 2024; Ben et al., 2025; Fu et al., 2024; Ze et al., 2025). OmniRetarget preserves human–object and human–environment interactions during retargeting, while SONIC tracks whole-body motion for physically grounded control (Yang et al., 2025a; Luo et al., 2026; Xie et al., 2025; Han et al., 2026). WholeBodyVLA learns loco-manipulation from vision and instructions, while OpenHLM and MotionWAM study whole-body humanoid policy and action-model training (Jiang et al., 2026; Hu et al., 2026; Zheng et al., 2026a; Xie et al., 2026). These approaches establish the control and learning foundation for humanoid loco-manipulation, but their supervision is collected or represented in a target-robot interface.

2.2 Egocentric Data and Cross-Embodiment Alignment

Egocentric interfaces and datasets provide a scalable source of visual and motion supervision. UMI records robot-compatible gripper trajectories and images, while EgoDex and EgoVerse provide large-scale RGB data with camera and hand-motion annotations (Chi et al., 2024; Hoque et al., 2026; Punamiya et al., 2026). Ψ_0 and EgoVLA use human data for embodied pretraining, and EMMA, EgoEngine, and Ego2Robot combine human video with robot learning or robot-data synthesis (Wei et al., 2026; Yang et al., 2025b; Zhu et al., 2026; Liu et al., 2026b; Wang et al., 2026). Related cross-embodiment methods address appearance and viewpoint transfer through cross-painting or robot-view augmentation (Chen et al., 2024a,b).

For humanoids, EgoHumanoid uses whole-body measurements, fixed-offset view alignment, and training-time perturbations (Shi et al., 2026). HALOMI adopts this view alignment and reference adaptation based on the cross-entropy method (CEM) for a dedicated controller (Zhao et al., 2026). KODA also uses command optimization, but derives onboard views from generated whole-body motion. Its canonical motion–observation records can be reused to augment manipulation datasets for compatible WBCs through controller-specific command generation and state recording.

3 System Framework

KODA converts egocentric human episodes into robot-view observations, canonical whole-body motion, and controller-compatible state–action data for VLA training as shown in Fig. 2. The conversion has three stages: Whole-Body Motion Generation in Sec. 3.2 produces a simulator-checked trajectory from sparse human references; Robot Onboard View Generation in Sec. 3.3 uses the corresponding body–camera trajectory to synthesize aligned observations, forming a controller-independent canonical record; and Control Command Generation in Sec. 3.4 tracks this motion with the selected controller to derive its state–action supervision. The canonical motion and observations are reused across interfaces, while controller-specific states and commands enable co-training with matching real-robot demonstrations. Qualitative demonstration timelines are provided in Appendix E.

3.1 Problem Setup

Let $\mathcal{R}$ denote the target humanoid and κ a deployment controller with its action interface. All sequences use the target control clock $t = 1, \dots, T$.

After geometry recovery, an RGB episode is represented as $e_h = (\{I_t^h, D_t^h, C_t^h, J_t^h\}_{t=1}^T, \ell)$, where I_t^h is the egocentric RGB image, D_t^h is observed or estimated depth, C_t^h is the source-camera pose, J_t^h contains the bilateral 3-D hand keypoints, and ℓ is the task instruction. Camera poses and hand keypoints come from source annotations or offline recovery. When source depth is unavailable, KODA estimates it from RGB using the metric-depth model of Depth Anything 3 (Lin et al., 2025). Depth, camera poses, and hand references must share a consistent metric coordinate system. Human episodes form $\mathcal{D}_{\text{human}}$. Real-robot episodes in $\mathcal{D}_{\text{robot}}^\kappa$ contain robot-view observations $I_t^\mathcal{R}$, controller-specific state features s_t^κ , actions a_t^κ , and instructions ℓ .

$\mathcal{C}_\mathcal{R}$ constructs a canonical record in $\mathcal{R}$ ’s observation and kinematic spaces, without controller-specific states or actions: $e_\mathcal{R}^{\text{can}} = \mathcal{C}_\mathcal{R}(e_h) = (\{I_t^\mathcal{R}, q_t^\mathcal{R}\}_{t=1}^T, \ell)$. Here $q_t^\mathcal{R}$ is the whole-body configuration corresponding to $I_t^\mathcal{R}$, including the floating-base pose and body and hand joint positions. The canonical trajectory provides a common motion reference for robot-view synthesis and controller-specific tracking.

The mappings $\mathcal{C}_{\mathcal{R}}$ and $\mathcal{G}_{\kappa}$ act episode-wise on datasets:

$$\begin{aligned} e_{\mathcal{R}}^{\kappa} &= \mathcal{G}_{\kappa}(e_{\mathcal{R}}^{\text{can}}) = \left(\{I_t^{\mathcal{R}}, s_t^{\kappa}, a_t^{\kappa}\}_{t=1}^T, \ell \right), \\ \mathcal{D}_{\text{KODA}}^{\kappa} &= \mathcal{G}_{\kappa}(\mathcal{C}_{\mathcal{R}}(\mathcal{D}_{\text{human}})), \quad \mathcal{D}_{\text{train}}^{\kappa} = \mathcal{D}_{\text{robot}}^{\kappa} \cup \mathcal{D}_{\text{KODA}}^{\kappa}. \end{aligned} \quad (1)$$

The mapping $\mathcal{G}_{\kappa}$ retains the canonical observations and instruction, adding state features and optimized command labels in κ 's interface as described in Sec. 3.4. Converted and real-robot episodes share a format and jointly train the VLA policy with parameters θ to predict $\hat{a}_t^{\kappa} = \pi_{\theta}(I_t^{\mathcal{R}}, s_t^{\kappa}, \ell)$.

3.2 Whole-Body Motion Generation

Given the camera poses and hand references in e_h , Whole-Body Motion Generation outputs the simulator-checked canonical whole-body motion $\{q_t^{\mathcal{R}}\}_{t=1}^T$.

Contact-Aware Hand Retargeting HAWOR (Zhang et al., 2025) estimates world-space MANO hand poses from egocentric RGB. Following AnyTeleop (Qin et al., 2023), KODA retargets them to obtain initial robot wrist poses and active finger-joint values using MANO palm geometry and a calibrated wrist-to-palm transform; distal mimic joints follow a fixed coupling map (Qin et al., 2023). Contact-aware refinement then adjusts this initialization on reliable contact frames. To identify these frames, EfficientSAM3 (Zeng et al., 2025) segments hand/forearm pixels, while nearby valid non-hand depths provide object-surface samples. Reliable contact frames are identified from multi-finger proximity, depth, and local surface consistency, then temporally filtered; without reliable contact, the initialization is retained.

At contact frames, wrist and finger poses are refined to preserve hand-object geometry. Following OmniRetarget (Yang et al., 2025a), selected human landmarks and corresponding robot-hand points are paired with the same nearby object samples to form connected world-frame interaction meshes. Each vertex descriptor is its offset from the uniformly weighted mean of its neighbors. We optimize wrist poses and finger joints to minimize differences between corresponding human and robot descriptors, preserving the demonstrated hand-object geometry. During optimization, object samples remain fixed, robot points are recomputed by forward kinematics, wrist translations, local rotations, and finger corrections are bounded, and deviations from initialization are penalized. Corrections are interpolated between contact frames and held constant outside their range. The optimized bilateral-wrist and finger trajectories are denoted by $w_{1:T}^{\text{ret}}$ and $f_{1:T}^{\text{ret}}$, respectively; the tracking stage uses C_t^h and $w_{1:T}^{\text{ret}}$ as sparse inputs. See Appendix A.1 for details.

Whole-body tracking via command optimization The source-camera poses are converted to torso poses using the camera-to-torso transform from the target robot's URDF, $p_t^{\text{torso}} = \mathcal{T}_{\text{cam} \rightarrow \text{torso}}(C_t^h)$. Combining p_t^{torso} with w_t^{ret} , concatenating the tracks in a fixed order, and applying a Savitzky-Golay filter (Schafer, 2011) yields $r_{1:T}^{\text{ref}}$.


Kimodo (Rempe et al., 2026) generates the initial whole-body configuration $q_{\text{init}}^{\mathcal{R}}$ from the first reference frame. After kinematic and physics checks, this configuration initializes the simulated rollout.

We add a three-point torso-and-bilateral-wrist tracking mode to ScaleBFM (Zeng et al., 2026) and fine-tune the pretrained whole-body controller $\phi_{\mathcal{R}}$ for the target embodiment. Its whole-body joint targets are applied by PD control in MuJoCo/MJX, but direct tracking can accumulate reference-tracking errors, as shown in TABLE 4. KODA therefore uses the cross-entropy method (CEM) (Pinneri et al., 2021) to optimize the sparse pose commands so that the realized motion tracks $r_{1:T}^{\text{ref}}$ more accurately while remaining simulation-feasible. Candidate commands are evaluated by simulated rollouts in a receding-horizon procedure detailed next and illustrated in Fig. 3; settings are reported in Sec. 4.1.

At each replan, $\mathbf{r}^{\text{ref}}$ spans H steps from simulated state x_{curr} . The latent vector $z \in \mathbb{R}^d$ contains 18 translation/rotation-vector coordinates per knot for torso and bilateral-wrist SE(3) residuals. CEM initializes diagonal-Gaussian parameters $(\mu, \sigma) = (0, \sigma_0)$ and samples one residual for each of the N parallel MJX environments per iteration:

$$\begin{aligned} z^{(n)} &= \mu + \sigma \odot \varepsilon^{(n)}, \quad \varepsilon_i^{(n)} \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, 1), \\ \tilde{\mathbf{r}}^{(n)} &= \text{Decode}(\mathbf{r}^{\text{ref}}, z^{(n)}), \end{aligned} \quad (2)$$

(a) Receding-horizon command correction



(b) CEM optimization within the planning horizon




Figure 3 Sampling-based command optimization (SCO) uses CEM to optimize sparse pose residuals through sampling, rollout, evaluation, and proposal updates. (a) Receding-horizon correction with a planning horizon and executed prefix. (b) The CEM loop proceeds from sampling to rollout, evaluation, and distribution update; the feedback arrow returns to sampling.

where $\odot$ denotes element-wise multiplication, $i = 1, \dots, d$ indexes latent coordinates, and $n = 1, \dots, N$ indexes candidates and environments. The operator `Decode` applies the sampled residuals at the reference knots, clips them to bounded physical ranges, and interpolates the translation and rotation corrections between knots to form the complete candidate pose-command trajectory $\hat{\mathbf{r}}^{(n)}$. We roll out the N candidates in parallel MJX environments:

$$\hat{\mathbf{q}}^{(n)} = \text{Rollout}_{\text{MJX}}(x_{\text{cur}}, \phi_{\mathcal{R}}, \hat{\mathbf{r}}^{(n)}), \quad \hat{\mathbf{r}}^{(n)} = \Pi_{\mathcal{B}}(\hat{\mathbf{q}}^{(n)}), \quad (3)$$

Here $\hat{\mathbf{q}}^{(n)}$ is the whole-body configuration trajectory realized by the n -th candidate command in MJX, and $\hat{\mathbf{r}}^{(n)} = \Pi_{\mathcal{B}}(\hat{\mathbf{q}}^{(n)})$ is the corresponding torso-and-wrist pose trajectory. The set $\mathcal{B}$ contains these torso and wrist bodies, and $\Pi_{\mathcal{B}}$ extracts their poses. Each realized trajectory is scored against the unperturbed reference by $J^{(n)} = d_{\mathcal{B}}(\hat{\mathbf{r}}^{(n)}, \mathbf{r}^{\text{ref}})$, a position/orientation tracking cost over these three links with receding-horizon time weights. Kinematic or physics violations are rejected. For the E valid lowest-cost samples indexed by $\mathcal{E}$, we update each dimension with $\mu_i^+ = \text{mean}_{n \in \mathcal{E}}[z_i^{(n)}]$ and $\sigma_i^+ = \max\{\sigma_{\min, i}, \text{std}_{n \in \mathcal{E}}[z_i^{(n)}]\}$, where the standard deviation uses divisor E and is bounded below by $\sigma_{\min, i}$. After I updates $(\mu, \sigma) \leftarrow (\mu^+, \sigma^+)$, we select the lowest-cost valid rollout among the nominal, best sampled, and final-mean candidates. Only the first Δ steps are retained; the resulting simulator state initializes the next replan, whose horizon advances by Δ . Concatenated prefixes form $q_{1:T}^{\mathcal{R}, \text{SCO}}$ for constrained refinement. See Appendix A.2 for implementation details.

Constrained Arm Motion Refinement To reduce residual wrist-tracking error, temporal Mink IK (Zakka, 2024) refines only the arm chain, including wrists. Contact-aware finger trajectories are retained; the optimized floating base, torso, and lower body remain fixed to preserve locomotion and balance motion. Refinement enforces joint limits, per-step displacement and velocity bounds, and self-collision constraints, while penalizing deviations from $q_{1:T}^{\mathcal{R}, \text{SCO}}$. Combining the refined arms and contact-aware fingers with the fixed non-arm



Figure 4 Qualitative example of the task “Put the cup into the drawer.” The rows show real-robot demonstrations (top), converted KODA data with aligned robot onboard observations (middle), and real-robot deployment of the learned VLA policy in the Robot-OOD setting (bottom). In the middle row, the main panels show the generated whole-body motion; the upper-left and upper-right insets show the source human view and the synthesized robot onboard view, respectively.

motion yields $q_{1:T}^{\mathcal{R}}$ for Robot Onboard View Generation. Replay shows comparable controller trackability after refinement, as quantified in Sec. 4.5. Further details are in Appendix A.3.

3.3 Robot Onboard View Generation

The optimized whole-body motion $q_{1:T}^{\mathcal{R}}$ from the previous section determines the robot camera trajectory. Because camera poses map camera coordinates to the world frame, the calibrated rigid torso–camera transform gives $C_t^{\mathcal{R}} = T_{\text{cam}}(q_t^{\mathcal{R}})$. KODA then synthesizes $I_t^{\mathcal{R}}$ from the source RGB frame I_t^h , associated depth D_t^h , source-camera pose C_t^h , and target configuration $q_t^{\mathcal{R}}$; all stages use this target camera pose and timestamp.

EfficientSAM3 (Zeng et al., 2025) predicts an arm/hand mask M_t^{arm} ; depth inside it is invalidated to obtain $\tilde{D}_t^h$. KODA back-projects remaining valid pixels using $\tilde{D}_t^h$, calibrated camera intrinsics, and C_t^h , transforms them to the target camera $C_t^{\mathcal{R}}$, and splats samples with a metric z-buffer, producing $(I_t^{\text{geo}}, D_t^{\text{geo}}, H_t^{\text{geo}})$: the reprojected image, depth map, and hole mask. MuJoCo renders $(I_t^{\text{rob}}, D_t^{\text{rob}})$ from the same $q_t^{\mathcal{R}}$ and camera. At pixels where both depth maps are valid, the nearer surface with smaller depth is retained: for pixel u , $M_t^{\text{rob}}(u) = 1$ if and only if $D_t^{\text{rob}}(u) < D_t^{\text{geo}}(u)$. If only one depth is valid, its corresponding pixel is retained; neighboring valid depths are used only inside geometric holes. The composite image is therefore $I_t^{\text{comp}} = M_t^{\text{rob}} I_t^{\text{rob}} + (1 - M_t^{\text{rob}}) I_t^{\text{geo}}$. DiffuEraser (Li et al., 2025) is applied only to residual holes $H_t^{\text{rem}} = H_t^{\text{geo}} \setminus M_t^{\text{rob}}$, namely, pixels left uncovered by both the reprojected scene and the rendered robot. Valid scene and robot pixels are frozen. The completed $I_t^{\mathcal{R}}$ is indexed on the target control clock and paired with $q_t^{\mathcal{R}}$. Together, the aligned observations $I_{1:T}^{\mathcal{R}}$, the whole-body motion $q_{1:T}^{\mathcal{R}}$ generated above, and the instruction ℓ form the canonical record $e_{\mathcal{R}}^{\text{can}}$, which the next section instantiates into controller-specific supervision. See Appendix B for details.

3.4 Control Command Generation and VLA Training

Given $e_{\mathcal{R}}^{\text{can}}$, KODA generates controller-specific command sequences for SONIC (Luo et al., 2026) and HEFT (Liu et al., 2026a) on the Unitree G1 with BrainCo hands. At 50 Hz, SCO optimizes a shared pelvis-pose and 29-body-joint residual space (hands excluded) against fixed $q_{1:T}^{\mathcal{R}}$. The G1 reference contains world-frame pelvis position and quaternion, 29 body-joint positions and velocities, and world-frame link poses from forward kinematics. Each candidate is decoded by the selected controller and rolled out in simulation with the same cost for both controllers—joint-position MSE, 14-link pose error, reference deformation, PD-target change, and terminal tracking penalties—and the same Control Command Generation settings reported in Sec. 4.1. The full search and objective are specified in Appendices C.1 and C.2.

Table 1 Real-robot task success (%) with increasing KODA data. Each controller uses 50 robot demonstrations per task. KODA demos denotes the number of converted human demonstrations per task; zero denotes robot-only training. Avg. is the macro-average over four tasks. Bold marks the best result within each controller, including ties.

Controller	KODA demos	Cup-to-basket		Parcel-to-shelf		Basket-to-shelf		Cup-to-drawer		Avg.	
		Robot-ID	Robot-OOD	Robot-ID	Robot-OOD	Robot-ID	Robot-OOD	Robot-ID	Robot-OOD	Robot-ID	Robot-OOD
SONIC	0	55.00	0.00	90.00	15.00	55.00	0.00	50.00	0.00	62.50	3.75
	25	55.00	10.00	100.00	25.00	50.00	15.00	70.00	15.00	68.75	16.25
	50	65.00	20.00	90.00	65.00	60.00	10.00	65.00	0.00	70.00	23.75
	75	70.00	35.00	90.00	70.00	65.00	30.00	60.00	30.00	71.25	41.25
	100	70.00	40.00	90.00	75.00	65.00	45.00	65.00	35.00	72.50	48.75
HEFT	0	65.00	0.00	85.00	20.00	65.00	0.00	60.00	0.00	68.75	5.00
	25	65.00	5.00	90.00	35.00	45.00	20.00	65.00	20.00	66.25	20.00
	50	60.00	30.00	80.00	55.00	65.00	25.00	70.00	15.00	68.75	31.25
	75	70.00	40.00	95.00	65.00	60.00	35.00	65.00	40.00	72.50	45.00
	100	80.00	50.00	90.00	75.00	70.00	40.00	65.00	40.00	76.25	51.25

The resulting actions a_t^κ use different controller interfaces. SONIC’s motion encoder maps the optimized motion references to 64-dimensional motion tokens (Luo et al., 2026). HEFT uses 29 absolute body-joint targets and a 9-dimensional pelvis code, giving 38 body-reference dimensions. The HEFT pelvis code represents roll r and pitch p as $(\sin r, \cos r - 1, \sin p, \cos p - 1)$, together with a wrapped next-step yaw change, three translation increments in the yaw-aligned frame, and absolute height. Planar position and heading accumulate from an anchor, whereas height is read directly. Both controllers append 12 hand-joint targets, giving 76-dimensional SONIC and 50-dimensional HEFT actions. Pose retargeting alone does not specify grasp force. KODA thresholds each retargeted robot hand’s mean finger curvature following EgoHumanoid (Shi et al., 2026). For grasp, hand-joint targets are set to a calibrated BrainCo grasp pose to encourage active closure; for release, targets equal the current retargeted joint angles. The 12 targets are sent separately to hand control.

Each converted episode records a 44-dimensional state vector s_t^κ in a format shared by both controllers: 29 body-joint angles, three-dimensional rollout gravity rotated to the robot base frame by the base quaternion, and 12 continuous BrainCo hand-joint angles from contact-aware retargeting as geometric hand-state estimates. Canonical images and the instruction are retained; TABLE 4 reports wrist/camera pose residuals, and TABLE 1 reports co-training utility. GR00T N1.7 predicts 40-step chunks of a_t^κ from $(I_t^\mathcal{R}, s_t^\kappa, \ell)$ using converted and real-robot episodes in the selected format. Changing the deployment interface reruns control command generation and state recording while reusing the canonical motion, observations, and instruction. Interface and export details are in Appendices C.3–C.5.

4 Evaluation

We evaluate whether KODA improves real-robot task success by augmenting fixed robot datasets with human demonstrations, reusing the same canonical records across SONIC and HEFT. EgoHumanoid (Shi et al., 2026) and HALOMI (Zhao et al., 2026) use different control interfaces and do not directly provide compatible whole-body state-action data. SONIC ablations assess the training benefits of observation alignment and controller-specific supervision; tracking diagnostics characterize motion fidelity and controller realization.

4.1 Experimental Setup

Platform, data, and evaluation. We use a Unitree G1 with BrainCo dexterous hands on four whole-body tasks: cup-to-basket, parcel-to-shelf, basket-to-shelf, and cup-to-drawer. Each task has 50 real-robot demonstrations from the corresponding SONIC or HEFT PICO teleoperation pipeline and 100 human egocentric demonstrations captured with a chest-mounted iPhone using MobileEgo Anywhere (Palanisamy et al., 2026). All 400 human clips include iPhone-recorded depth. Both sources use 50 Hz. The scaling study evaluates SONIC (Luo et al., 2026) and HEFT (Liu et al., 2026a); component ablations use SONIC with 50 robot and 100 converted-human demonstrations per task. Per controller configuration, GR00T N1.7 (Bjorck et al., 2025) trains on all four tasks for 20,000 iterations with batch size 160. Within each task, robot/human sampling is proportional to episode counts; other sampling settings follow GR00T defaults. We report one training run per configuration and 20 real-robot trials per task and condition. This trial count matches EgoHumanoid’s 20 trials per setting (Shi et al., 2026) and exceeds the 10 trials per task and model reported in Ψ_0 (Wei et al., 2026). Success requires

Table 2 Effect of observation alignment on SONIC policy success (%). The four variants use the same demonstrations, state-action supervision, policy architecture, training budget, and evaluation protocol. Robot-ID and Robot-OOD are reported separately for each variant. Bold marks the highest success rate for each task and condition, including ties.

Task	No alignment		Fixed offset		Motion-based reprojection		Full alignment	
	Robot-ID	Robot-OOD	Robot-ID	Robot-OOD	Robot-ID	Robot-OOD	Robot-ID	Robot-OOD
Cup-to-basket	50	15	60	15	65	25	70	40
Parcel-to-shelf	85	35	80	45	85	65	90	75
Basket-to-shelf	65	20	60	25	60	35	65	45
Cup-to-drawer	50	5	55	10	55	25	65	35
Macro-average	62.5	18.75	63.75	23.75	66.25	37.5	72.5	48.75

completing the placement without a fall or unrecovered manipulation failure. Reported averages weight the four tasks equally.

Conditions covered by human demonstrations. Robot-ID uses object appearance and pose conditions present in the robot demonstrations. Robot-OOD uses conditions absent from those demonstrations: a green bottle is replaced by a brown bottle in cup-to-basket and cup-to-drawer, a purple parcel by a white parcel in parcel-to-shelf, and the basket handle is tilted by approximately 30° in basket-to-shelf. The human training pool contains 80 shifted and 20 robot-matched clips per task. Robot-OOD tests whether human demonstrations supply coverage missing from the robot dataset, rather than zero-shot generalization to conditions absent from all training data. Data and evaluation details are in Appendices D.1 and D.2.

CEM settings. Both stages run at 50 Hz and use CEM with 80 candidates per iteration, 8 iterations per replan, and a 40-step horizon. Whole-Body Motion Generation uses 10 elites and executes 10 steps before replanning. Control Command Generation with SONIC and HEFT uses 9 elites and executes 20 steps before replanning.

4.2 Extending Robot Demonstration Coverage

Each task combines 50 fixed robot demonstrations with 0, 25, 50, 75, or 100 converted human demonstrations, randomly sampled independently at each budget from its 100-clip pool without nesting. Policy architecture, training budget, and evaluation protocol remain fixed. TABLE 1 shows that KODA provides its largest gains on conditions missing from robot demonstrations. With 100 converted demonstrations, SONIC Robot-OOD success rises from 3.75% to 48.75% (+45 percentage points), and HEFT rises from 5.00% to 51.25% (+46.25 points). Robot-ID success also improves, from 62.5% to 72.5% for SONIC and from 68.75% to 76.25% for HEFT. Robot-OOD macro-average trends upward over tested budgets for both controllers. Individual-task and Robot-ID results vary at intermediate data sizes, so we focus on overall gains at the full data budget. KODA thus supplies useful supervision for additional conditions covered by human demonstrations without adding robot demonstrations. Fig. 4 illustrates converted data and real-robot execution in the Robot-OOD setting for cup-to-drawer. See Appendix F for the other three tasks. The same converted human demonstrations yield large Robot-OOD gains under both SONIC and HEFT, demonstrating cross-controller reuse, as shown in TABLE 1. The following ablations isolate observation alignment and control command generation.

4.3 Observation Alignment

We hold the converted demonstrations, state-action supervision, policy architecture, training budget, and evaluation protocol fixed across four tasks. No alignment retains the original human images. Fixed offset uses a 0.10 m vertical camera shift with ± 0.05 m training-time perturbations, adapting the fixed-offset strategy of (Shi et al., 2026) to chest-mounted capture. Fixed offset and Motion-based reprojection use only reprojection and hole inpainting; Motion-based reprojection derives the target camera poses from the canonical robot motion. Full alignment additionally removes human arms and renders and composites the target robot. See Appendix D.3 for details.

Full alignment reaches 48.75% Robot-OOD success versus 23.75% with the adapted EgoHumanoid fixed-

Table 3 Effect of controller-specific command optimization on SONIC policy success (%). Both variants use 50 robot and 100 converted-human demonstrations per task, with identical canonical trajectories and synthesized observations. Without SCO omits optimization only during Control Command Generation; With SCO uses optimized commands. Policy architecture, training budget, and evaluation protocol are held fixed. Bold marks the higher success rate for each task and condition.

Task	Robot-ID		Robot-OOD	
	Without SCO	With SCO	Without SCO	With SCO
Cup-to-basket	25	70	0	40
Parcel-to-shelf	70	90	10	75
Basket-to-shelf	50	65	0	45
Cup-to-drawer	50	65	20	35
Macro-average	48.75	72.5	7.5	48.75

SCO: sampling-based command optimization. Macro-average weights the four tasks equally.

offset baseline (Shi et al., 2026) and 18.75% without alignment (TABLE 2). Motion-based reprojection reaches 37.5%. Fixed offset, Motion-based reprojection, and Full alignment yield Robot-ID point estimates of 63.75%/66.25%/72.5%, respectively.

4.4 Controller-Consistent State-Action Supervision

The variant without SCO directly uses the generated canonical whole-body motion as the tracking controller’s reference, omitting optimization during control command generation while retaining Whole-Body Motion Generation and Robot Onboard View Generation. Both variants use identical canonical trajectories, observations, demonstrations, policy architecture, training budget, and evaluation protocol. Since body states come from controller rollouts, the comparison evaluates state-action supervision as a whole. See Appendix D.3 for details.

TABLE 3 shows that optimized control command generation raises macro-average success from 48.75% to 72.5% Robot-ID and from 7.5% to 48.75% Robot-OOD, gains of 23.75 and 41.25 percentage points. Both metrics improve on all four tasks. Thus, controller-command generation also affects real-robot Robot-ID performance. Without this optimization, Robot-ID success falls below the 62.5% robot-only baseline in TABLE 1; the complete pipeline reaches 72.5%. Controller-consistent supervision thus matters to whether the added data benefits the policy. Together, the ablations support aligning both observations and state-action supervision with deployment.

4.5 Motion Fidelity and Controller Trackability

Motion-generation entries in TABLE 4 average aggregate errors from LAFAN (Harvey et al., 2020) and 400 egocentric training clips. LAFAN provides a check on highly dynamic motions; both controllers share evaluation data.

Wrist errors decrease to 3.03 cm/0.1106 rad with SCO and 0.33 cm/0.0296 rad after Mink refinement. Controller-specific SCO improves wrist/camera tracking except SONIC camera orientation (+0.0024 rad). Replaying the motions before and after Mink refinement with the same SONIC tracker on both sources yields comparable errors: pelvis-relative non-arm/wrist positions 2.60/1.59 → 2.73/1.75 cm, torso orientation 0.0588 → 0.0579 rad, and non-arm joint RMSE 0.0750 → 0.0784 rad. Metric and replay details are in Appendix D.4. TABLES 2 and 3 show task benefits from observation alignment and controller-specific supervision. TABLE 1 supports the reusable records’ training utility across both controllers despite residual tracking mismatch.

Table 4 Tracking errors for whole-body motion generation and controller rollouts. Position and orientation are reported separately in centimeters and radians; lower is better. Generation errors use the source motion as the reference, while rollout errors use the fixed canonical motion. Bold marks the smallest reported value for each target and error type, including ties.

Stage	Target	Direct		SCO		SCO+IK	
		Position (cm)	Orientation (rad)	Position (cm)	Orientation (rad)	Position (cm)	Orientation (rad)
Whole-body motion generation	Wrist	15.02	0.2046	3.03	0.1106	0.33	0.0296
	Torso	–	0.1302	–	0.0930	–	0.0930
SONIC rollout	Wrist	61.50	0.1777	2.88	0.1292	–	–
	Camera	60.90	0.0784	3.20	0.0808	–	–
HEFT rollout	Wrist	19.61	0.0677	2.94	0.0552	–	–
	Camera	19.51	0.0426	2.84	0.0337	–	–

Direct uses unoptimized references. SCO uses CEM-optimized commands; SCO+IK adds Mink arm refinement during whole-body motion generation. A dash (–) denotes an unreported metric or variant.

5 Limitations

KODA does not yet fully reproduce fine-grained hand–object interactions, limiting their fidelity and physical feasibility. In the future, we plan to incorporate real-to-sim scene reconstruction and contact modeling to better recover object geometry and interaction dynamics. We will also broaden human data sources beyond egocentric recordings, including third-person videos and motion capture, converting diverse human demonstrations into canonical whole-body records for controller-specific state–action generation and policy training.

6 Conclusion

KODA converts egocentric human demonstrations into canonical whole-body motion and onboard observations, then generates controller-specific state–action supervision for VLA training. Reusing these records with fixed robot data improves VLA task success under SONIC and HEFT on four G1 tasks, especially on conditions covered only by human demonstrations. Reduced tracking errors and co-training gains despite residual mismatch support expanding robot-policy coverage.

References

- Qingwei Ben, Feiyu Jia, Jia Zeng, Junting Dong, Dahua Lin, and Jiangmiao Pang. Homie: Humanoid loco-manipulation with isomorphic exoskeleton cockpit. *arXiv preprint arXiv:2502.13013*, 2025.
- Johan Bjorck, Fernando Castañeda, Nikita Cherniadev, Xingye Da, Runyu Ding, Linxi Fan, Yu Fang, Dieter Fox, Fengyuan Hu, Spencer Huang, et al. Gr00t n1: An open foundation model for generalist humanoid robots. *arXiv preprint arXiv:2503.14734*, 2025.
- Lawrence Yunliang Chen, Kush Hari, Karthik Dharmarajan, Chenfeng Xu, Quan Vuong, and Ken Goldberg. Mirage: Cross-embodiment zero-shot policy transfer with cross-painting. *arXiv preprint arXiv:2402.19249*, 2024a.
- Lawrence Yunliang Chen, Chenfeng Xu, Karthik Dharmarajan, Muhammad Zubair Irshad, Richard Cheng, Kurt Keutzer, Masayoshi Tomizuka, Quan Vuong, and Ken Goldberg. Rovi-aug: Robot and viewpoint augmentation for cross-embodiment robot learning. *arXiv preprint arXiv:2409.03403*, 2024b.
- Cheng Chi, Zhenjia Xu, Chuer Pan, Eric Cousineau, Benjamin Burchfiel, Siyuan Feng, Russ Tedrake, and Shuran Song. Universal manipulation interface: In-the-wild robot teaching without in-the-wild robots. *arXiv preprint arXiv:2402.10329*, 2024.
- Zipeng Fu, Qingqing Zhao, Qi Wu, Gordon Wetzstein, and Chelsea Finn. Humanplus: Humanoid shadowing and imitation from humans. *arXiv preprint arXiv:2406.10454*, 2024.
- Jinrui Han, Weiwei Xie, Jiakun Zheng, Jiyuan Shi, Weinan Zhang, Ting Xiao, and Chenjia Bai. KungfuBot2: Learning versatile motion skills for humanoid whole-body control. In *2026 IEEE International Conference on Robotics and Automation (ICRA)*, 2026.
- Félix G Harvey, Mike Yurick, Derek Nowrouzezahrai, and Christopher Pal. Robust motion in-betweening. *ACM Trans. Graph.*, 39(4):60:1–60:12, 2020. doi: 10.1145/3386569.3392480.
- Tairan He, Zhengyi Luo, Xialin He, Wenli Xiao, Chong Zhang, Weinan Zhang, Kris Kitani, Changliu Liu, and Guanya Shi. Omnih2o: Universal and dexterous human-to-humanoid whole-body teleoperation and learning. *arXiv preprint arXiv:2406.08858*, 2024.
- Ryan Hoque, Peide Huang, David J. Yoon, Mouli Sivapurapu, and Jian Zhang. EgoDex: Learning dexterous manipulation from large-scale egocentric video. In *Proc. Int. Conf. Learn. Represent. (ICLR)*, volume 2026, pages 4218–4237, 2026.
- Yingdong Hu, Haodong Zhu, Boyuan Zheng, Yihang Hu, Tong Zhang, Zunhao Chen, Junming Zhao, Ruiqian Nai, and Yang Gao. Openhlm: An empirical recipe for whole-body humanoid loco-manipulation. *arXiv preprint arXiv:2606.22174*, 2026.
- Haoran Jiang, Jin Chen, Qingwen Bu, Li Chen, Modi Shi, Yanjie Zhang, DeLong Li, Chuanzhe Suo, Chuang Wang, Zhihui Peng, and Hongyang Li. WholeBodyVLA: Towards unified latent VLA for whole-body loco-manipulation control. In *Proc. Int. Conf. Learn. Represent. (ICLR)*, volume 2026, pages 157438–157461, 2026.
- Xiaowen Li, Haolan Xue, Peiran Ren, and Liefeng Bo. Diffuseraser: A diffusion model for video inpainting. *arXiv preprint arXiv:2501.10018*, 2025.
- Haotong Lin, Sili Chen, Jun Hao Liew, Donny Y. Chen, Zhenyu Li, Guang Shi, Jiashi Feng, and Bingyi Kang. Depth Anything 3: Recovering the visual space from any views. *arXiv preprint arXiv:2511.10647*, 2025.
- Chenxin Liu, Qingzhou Lu, Guangxiao Yang, Xuanyang Shi, Chenghan Yang, Yanjiang Guo, and Jianyu Chen. HEFT: Heavy-payload full-size humanoid teleoperation with privileged motion guidance and windowed payload curriculum. *arXiv preprint arXiv:2607.02332*, 2026a.
- Yangcen Liu, Shuo Cheng, Xinchun Yin, Woo Chul Shin, Alfred Cueva, Yiran Yang, Zhenyang Chen, Chuyue Zhang, and Danfei Xu. Egoengine: From egocentric human videos to high-fidelity dexterous robot demonstrations. *arXiv preprint arXiv:2606.12604*, 2026b.
- Zhengyi Luo, Ye Yuan, Tingwu Wang, Chenran Li, Fernando Castañeda, Sirui Chen, Zi-Ang Cao, Jiefeng Li, David Minor, Qingwei Ben, et al. Sonic: Supersizing motion tracking for natural humanoid whole-body control. *Science Robotics*, 11(117):eaed4592, 2026.
- Senthil Palanisamy, Abhishek Anand, Satpal Singh Rathore, Pratyush Patnaik, Shubhanshu Khatana, and Ekaksh Janweja. Mobileego anywhere: Open infrastructure for long horizon egocentric data on commodity hardware. *arXiv preprint arXiv:2605.05945*, 2026.

-
- Cristina Pinneri, Shambhuraj Sawant, Sebastian Blaes, Jan Achterhold, Joerg Stueckler, Michal Rolinek, and Georg Martius. Sample-efficient cross-entropy method for real-time planning. In *Conference on Robot Learning*, pages 1049–1065. PMLR, 2021.
- Ryan Punamiya, Simar Kareer, Zeyi Liu, Josh Citron, Ri-Zhao Qiu, Xiongyi Cai, Alexey Gavryushin, Jiaqi Chen, Davide Liconti, Lawrence Y Zhu, et al. Egoverse: An egocentric human dataset for robot learning from around the world. *arXiv preprint arXiv:2604.07607*, 2026.
- Yuzhe Qin, Wei Yang, Binghao Huang, Karl Van Wyk, Hao Su, Xiaolong Wang, Yu-Wei Chao, and Dieter Fox. Anyteleop: A general vision-based dexterous robot arm-hand teleoperation system. *arXiv preprint arXiv:2307.04577*, 2023.
- Davis Rempe, Mathis Petrovich, Ye Yuan, Haotian Zhang, Xue Bin Peng, Yifeng Jiang, Tingwu Wang, Umar Iqbal, David Minor, Michael de Ruyter, et al. Kimodo: Scaling controllable human motion generation. *arXiv preprint arXiv:2603.15546*, 2026.
- Ronald W Schafer. What is a savitzky-golay filter?[lecture notes]. *IEEE Signal processing magazine*, 28(4):111–117, 2011.
- Modi Shi, Shijia Peng, Jin Chen, Haoran Jiang, Tianyu Li, Di Huang, Ping Luo, Hongyang Li, and Li Chen. Egohumanoid: Unlocking in-the-wild loco-manipulation with robot-free egocentric demonstration. *arXiv preprint arXiv:2602.10106*, 2026.
- Ye Wang, Pei Lin, Xiong-Hui Chen, Haoqi Yuan, Zhixuan Liang, Yiyang Huang, Anzhe Chen, Zixing Lei, Jie Zhang, Tao Zhang, et al. Ego2robot: Scalable robot data synthesis from egocentric human data. *arXiv preprint arXiv:2608.02580*, 2026.
- Songlin Wei, Hongyi Jing, Boqian Li, Zhenyu Zhao, Jiageng Mao, Zhenhao Ni, Sicheng He, Jie Liu, Xiawei Liu, Kaidi Kang, et al. Ψ_0 : An open foundation model towards universal humanoid loco-manipulation. *arXiv preprint arXiv:2603.12263*, 2026.
- Weiji Xie, Jinrui Han, Jiakun Zheng, Huanyu Li, Xinzhe Liu, Jiyuan Shi, Weinan Zhang, Chenjia Bai, and Xuelong Li. KungfuBot: Physics-based humanoid whole-body control for learning highly-dynamic skills. In *Advances in Neural Information Processing Systems*, 2025.
- Weiji Xie, Jiakun Zheng, Jinrui Han, Jiyuan Shi, Weinan Zhang, Chenjia Bai, and Xuelong Li. TextOp: Real-time interactive text-driven humanoid robot motion generation and control. In *Conference on Robot Learning*, 2026.
- Lujie Yang, Xiaoyu Huang, Zhen Wu, Angjoo Kanazawa, Pieter Abbeel, Carmelo Sferrazza, C Karen Liu, Rocky Duan, and Guanya Shi. Omniretarget: Interaction-preserving data generation for humanoid whole-body loco-manipulation and scene interaction. *arXiv preprint arXiv:2509.26633*, 2025a.
- Ruihan Yang, Qinxu Yu, Yecheng Wu, Rui Yan, Borui Li, An-Chieh Cheng, Xueyan Zou, Yunhao Fang, Xuxin Cheng, Ri-Zhao Qiu, et al. Egovla: Learning vision-language-action models from egocentric human videos. *arXiv preprint arXiv:2507.12440*, 2025b.
- Kevin Zakka. Mink: Python inverse kinematics based on mujoco, 2024. URL <https://github.com/kevinzakka/mink>.
- Yanjie Ze, Siheng Zhao, Weizhuo Wang, Angjoo Kanazawa, Rocky Duan, Pieter Abbeel, Guanya Shi, Jiajun Wu, and C Karen Liu. Twist2: Scalable, portable, and holistic humanoid data collection system. *arXiv preprint arXiv:2511.02832*, 2025.
- Chengxi Zeng, Yuxuan Jiang, and Aaron Zhang. EfficientSAM3: Progressive hierarchical distillation for video concept segmentation from sam1, 2, and 3. *arXiv preprint arXiv:2511.15833*, 2025.
- Weishuai Zeng, Kangning Yin, Xiaojie Niu, Shunlin Lu, Weixiang Zhong, Jiahe Chen, Feiyu Jia, Xiao Chen, Zirui Wang, Furui Xu, et al. Scaling behavior foundation model for humanoid robots. *arXiv preprint arXiv:2607.15163*, 2026.
- Jinglei Zhang, Jiankang Deng, Chao Ma, and Rolandos Alexandros Potamias. Hawor: World-space hand motion reconstruction from egocentric videos. In *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 1805–1815. IEEE, 2025.
- Zehui Zhao, Yuxuan Zhao, Gaojing Zhang, Chenxi Liu, Maolin Zheng, and Wenzhao Lian. Halomi: Learning humanoid loco-manipulation with active perception from human demonstrations. *arXiv preprint arXiv:2606.18772*, 2026.
- Jia Zheng, Teli Ma, Yudong Fan, Zifan Wang, Shuo Yang, and Junwei Liang. Motionwam: Towards foundation world action models for real-time humanoid loco-manipulation. *arXiv preprint arXiv:2606.09215*, 2026a.

Ruijie Zheng, Dantong Niu, Yuqi Xie, Jing Wang, Mengda Xu, Yunfan Jiang, Fernando Castañeda, Fengyuan Hu, You Liang Tan, Letian Fu, et al. Egoscale: Scaling dexterous manipulation with diverse egocentric human data. *arXiv preprint arXiv:2602.16710*, 2026b.

Lawrence Y Zhu, Pranav Kuppili, Ryan Punamiya, Patcharapong Aphiwetsa, Dhruv Patel, Simar Kareer, Sehoon Ha, and Danfei Xu. EMMA: Scaling mobile manipulation via egocentric human data. *IEEE Robot. Autom. Lett.*, 11(3): 3087–3094, March 2026. doi: 10.1109/LRA.2026.3653320.

Appendix

Appendix A expands contact-aware retargeting, sparse command optimization, and arm refinement. Appendix B specifies coordinate and visibility conventions. Appendix C details controller-specific search, action encoding, and state construction. Appendix D clarifies experimental conditions, comparisons, and metric references. Appendix E presents complete-clip demonstration timelines. Appendix F provides Robot-ID, KODA, and Robot-OOD comparisons for the remaining three tasks.

A Whole-Body Motion Generation Details

This appendix expands Sec. 3.2. Source-camera and bilateral hand references share a metric world frame. Retargeted wrist and finger tracks $w_{1:T}^{\text{ret}}, f_{1:T}^{\text{ret}}$ supply the hand references; $r_{1:T}^{\text{ref}}$ denotes the desired torso and wrist poses. SCO produces realized whole-body motion $q_{1:T}^{\mathcal{R}, \text{SCO}}$. Arm refinement yields the canonical floating-base, body, and hand sequence $q_{1:T}^{\mathcal{R}}$ while retaining the hand references and non-arm motion.

A.1 Contact-aware hand retargeting

Initialization and contact evidence. HAWOR (Zhang et al., 2025) supplies world-space MANO hand estimates. Retargeting following AnyTeleop (Qin et al., 2023) uses palm geometry and a calibrated wrist-to-palm transform to initialize each robot wrist and its active finger joints. Distal mimic joints are determined by the coupling map, so they are not independent optimization variables. A correspondence between selected human landmarks and robot-hand points provides the geometric basis for refinement. The correspondences express analogous locations on the two hands; they do not require their shapes or articulation ranges to coincide.

The hand/forearm segmentation removes human-hand pixels from the depth samples used as object evidence. Nearby valid non-hand depths provide local object surface points. Multi-finger proximity, depth validity, and local surface consistency identify candidate contact frames, which are temporally filtered. These checks determine where geometric refinement is supported by the observation. When an entire sequence has no reliable contact, the initialized hand trajectory is retained. Contact evidence is used to recover spatial hand-object relationships; it is not a measurement of contact force.

Local wrist and finger corrections. At one contact frame, let (p^0, R^0) be the initialized wrist pose and f^0 the active finger-joint vector for one hand. In the implemented six-active-joint hand model, $\delta p, \delta \omega \in \mathbb{R}^3$ and $\delta f \in \mathbb{R}^6$. We optimize $\xi = (\delta p, \delta \omega, \delta f) \in \mathbb{R}^{12}$ using

$$\begin{aligned} p(\xi) &= p^0 + R^0 \delta p, \\ R(\xi) &= R^0 \text{Exp}([\delta \omega]_{\times}), \\ f(\xi) &= \text{clip}(f^0 + \delta f, f_{\min}, f_{\max}). \end{aligned} \tag{A.1}$$

The initialized position p^0 is in the world frame, and R^0 maps wrist coordinates to world coordinates. Both wrist corrections are expressed relative to the initialized wrist frame. Translations are in meters; rotation vectors and finger-joint angles are in radians. The clip is componentwise, and $f_{\min}, f_{\max} \in \mathbb{R}^6$ are the active-joint limits in the robot’s MJCF model. Here $[\delta \omega]_{\times}$ is the skew-symmetric matrix associated with the rotation vector, and Exp maps it to a valid rotation. Forward kinematics evaluates the robot landmarks under the corrected wrist and finger configuration, including the fixed mimic-joint relationships. Thus wrist placement and finger articulation affect the same geometric objective.

Shared interaction geometry. Let P^h be the selected human landmarks, $P^{\mathcal{R}}(\xi)$ their robot counterparts, and O the object samples for this frame. Following the interaction-mesh construction of OmniRetarget (Yang et al., 2025a), the ordered vertex arrays $V^h = [P^h; O]$ and $V^{\mathcal{R}}(\xi) = [P^{\mathcal{R}}(\xi); O]$ share vertex correspondences and connectivity. The implementation constructs adjacency from a Delaunay tetrahedralization of the human landmarks and object samples, with a nearest-neighbor fallback when tetrahedralization fails. Connectivity is then held fixed throughout the frame’s optimization and reused for the robot vertices.

Let $v_i \in \mathbb{R}^3$ be the world-frame position of the i th vertex of V . For each vertex with a nonempty neighbor set $\mathcal{N}(i)$, its descriptor is

$$L_i(V) = v_i - \frac{1}{|\mathcal{N}(i)|} \sum_{j \in \mathcal{N}(i)} v_j. \quad (\text{A.2})$$

The uniform weights sum to one. Consequently, adding a common translation to the vertices leaves their descriptors unchanged. Hand-object relationships still constrain the wrist because the observed object vertices remain fixed while the robot-hand vertices move. Object descriptors can also change when their neighbors include hand vertices, even though the object positions themselves are constant. The descriptor discrepancy therefore includes both hand and object rows:

$$\xi^* = \arg \min_{\xi \in \mathcal{A}} \lambda_L \sum_{i \in \mathcal{V}} \|L_i(V^{\mathcal{R}}(\xi)) - L_i(V^h)\|_2^2 + \mathcal{L}_{\text{reg}}(\xi). \quad (\text{A.3})$$

Here $\mathcal{V}$ indexes corresponding vertices and λ_L weights descriptor matching. Let $\xi^- = (\delta p^-, \delta \omega^-, \delta f^-)$ denote the preceding solved contact-frame correction, with $\xi^- = 0$ for the first contact frame, and define $[a]_+ = \max(a, 0)$. The regularizer is

$$\begin{aligned} \mathcal{L}_{\text{reg}}(\xi) = & \alpha_w^2 (\|\delta p\|_2^2 + \|\delta \omega\|_2^2) + \alpha_f^2 \|\delta f\|_2^2 \\ & + \beta_w^2 (\|\delta p - \delta p^-\|_2^2 + \|\delta \omega - \delta \omega^-\|_2^2) \\ & + \beta_f^2 \|\delta f - \delta f^-\|_2^2 \\ & + \gamma^2 ([\|\delta p\|_2 - b_t]_+^2 + [\|\delta \omega\|_2 - r_{\text{max}}]_+^2). \end{aligned}$$

The coefficients α_w, α_f weight wrist and finger anchoring; β_w, β_f weight temporal changes in those corrections; and γ weights the soft magnitude limits. These coefficients multiply residuals, so their squares appear in the objective. The first two terms anchor the correction to its initialization, the next two penalize changes from the preceding contact solution, and the last term softly limits translation and rotation magnitudes. We use the following default weights:

Term	Parameter	Value
Laplacian descriptor loss	λ_L	10
Wrist anchor residual	α_w	0.10
Finger anchor residual	α_f	0.018
Wrist temporal residual	β_w	12
Finger temporal residual	β_f	2
Soft-limit residual	γ	40

The correction box at frame t is

$$\mathcal{A} = \{\xi : |\delta p_i| \leq b_t, \quad |\delta \omega_i| \leq r_{\text{max}} \quad (i = 1, 2, 3), \\ |\delta f_j| \leq \delta f_{\text{max}} \quad (j = 1, \dots, 6)\}.$$

Here b_t is the frame's translation bound in meters, r_{max} is the rotation-vector coordinate bound in radians, and δf_{max} is the finger-correction coordinate bound in radians. The default angular bounds are $r_{\text{max}} = \pi/15$ rad (12°) and $\delta f_{\text{max}} = 0.28$ rad. These are componentwise box bounds; the rotation-vector norm also enters the soft magnitude penalty above. Active-joint limits are applied separately by the clipping in Eq. (A.1). Let $d_{0,t}$ be the world-frame distance from the initialized palm center to the centroid of the object samples used in the Laplacian graph, with $d_{0,t} = 0$ for an empty sample set. The default translation bound is

$$b_t = \begin{cases} \min(d_{0,t}, 0.040 \text{ m}), & d_{0,t} \text{ is finite and } d_{0,t} \geq 0.025 \text{ m}, \\ 0.015 \text{ m}, & \text{otherwise.} \end{cases}$$

These coordinate bounds and soft magnitude penalties restrict displacement when local depth evidence is incomplete.

Numerical refinement and temporal propagation. The contact solver stacks the descriptor, anchor, temporal, and soft-limit residuals into a vector $\rho(\xi)$ so that $\|\rho(\xi)\|_2^2$ equals the objective in Eq. (A.3). In particular, descriptor differences are multiplied by $\sqrt{\lambda_L}$; the remaining residual blocks use $\alpha_w, \alpha_f, \beta_w, \beta_f, \gamma$ as defined above. Let J_ρ denote the Jacobian of this weighted residual vector with respect to ξ . The fixed Laplacian matrix L satisfies $(LV)_i = L_i(V)$, and I_3 is the 3×3 identity. The unweighted descriptor Jacobian follows by applying $L \otimes I_3$ to the robot-vertex Jacobian; object-vertex derivatives are zero. Translation derivatives follow the wrist transform analytically. Rotation and finger-joint derivatives use forward finite differences with an increment of 10^{-4} rad. Each iteration computes a damped linearized step,

$$\delta\xi = -(J_\rho^\top J_\rho + \lambda_d I)^{-1} J_\rho^\top \rho, \quad (\text{A.4})$$

where I is the 12×12 identity and λ_d is the damping coefficient, set to 10^{-4} . The solver first clips the step against the remaining coordinate bounds, then rescales it if its Euclidean norm exceeds the default trust-region radius of 0.2. This norm uses the correction coordinates in meters and radians without additional coordinate normalization. The corrected pose is re-evaluated before the next iteration. The default iteration limits are 20 for the first contact frame and 10 for subsequent contact frames. Iterations stop earlier when successive squared-residual costs satisfy NumPy’s default `isclose` test or when the clipped and rescaled step has norm below 2×10^{-4} ; in the latter case, the small step is applied before termination. The preceding contact solution initializes the next contact-frame solve.

Within the contact optimizer, correction candidates are propagated to the target clock by linear interpolation of translation and finger offsets and shortest-path interpolation of the correction rotations on $SO(3)$ between contact intervals. Before the first contact, the correction ramps from zero at the sequence start to the first contact correction; after the last contact it ramps to zero at the sequence end. These ramps apply when the corresponding non-contact frames exist. With no contact frames, all corrections are zero. The propagated correction coordinates are then smoothed with a separate Savitzky–Golay filter whose default window contains nine frames and whose polynomial degree is two, using polynomial edge evaluation (`interp` mode). Sequences shorter than nine frames are left unchanged by this filter. Filtering does not impose the prefilter endpoint values as equality constraints. This correction filter is distinct from the sparse-reference filter described in Sec. A.3. The resulting $f_{1:T}^{\text{ret}}$ is retained as a continuous geometric hand reference; the grasp and release command construction is a separate operation described in Sec. 3.4.

For the tracking pipeline, the wrist reference is assembled from the candidate wrist poses over the interval between the first and last contact frames. Outside this interval, the first-contact residual is held before contact, and the last-contact residual is held after contact. The held quantities are the world-frame translation offset and the local rotation offset composed on the right of the nominal wrist orientation. They are applied relative to the nominal wrist pose at each frame, so holding these residuals does not fix the absolute wrist pose. This boundary rule applies to the injected wrist reference and does not replace the interpolated and smoothed finger trajectory. The assembled three-point reference is subsequently filtered as described in Sec. A.3.

A.2 Sparse pose-command optimization

Search space and decoding. The source-camera poses are mapped to torso poses using the target robot’s URDF camera-to-torso transform. Together with the retargeted wrists, these produce three ordered reference tracks: torso, left wrist, and right wrist. The ScaleBFM-based tracker (Zeng et al., 2026) receives these sparse pose commands and produces whole-body PD targets. SCO changes the sparse commands and evaluates their realized motion; it does not directly search all robot joint configurations.

At each replanning instant, $K = 10$ knots span a horizon of $H = 40$ control steps. For each tracked body b and knot k , the latent contains translation and rotation-vector residuals $(z_{kb}^p, z_{kb}^R) \in \mathbb{R}^6$. The three bodies therefore contribute 18 coordinates per knot and 180 coordinates in total. Normalized coordinates are clipped to $[-1, 1]$ before physical scaling:

$$\delta p_{kb} = 0.055 \text{ clip}(z_{kb}^p, -1, 1), \quad \delta \omega_{kb} = 0.0825 \text{ clip}(z_{kb}^R, -1, 1). \quad (\text{A.5})$$

Translation is in meters and rotation is in radians. The decoder interpolates corrections between knots, applies translation offsets to the nominal positions, and composes rotations represented by the rotation vectors with the nominal orientations. Joint references, pelvis guidance, and other body references are not independently perturbed by this search. Their realized motion is determined by the tracker and simulator.

Table A.1 SCO settings for Whole-Body Motion Generation. The residual and proposal scales apply to this stage’s three-point pose-command space.

Quantity	Value
Control frequency	50 Hz
Planning horizon H / executed prefix Δ	40 / 10 steps
Temporal knots K / latent dimension d	10 / 180
Candidates N / CEM updates I / elites E	80 / 8 / 10
Initial standard deviation / floor	1.6 / 0.075
Translation / rotation scale	0.055 m / 0.0825 rad
Position / orientation cost weight	1.0 / 0.1
Torso / left wrist / right wrist weight	0.25 / 1.0 / 1.0

Tracking objective. All candidates start from the same current simulator state and are scored against the unmodified reference $\mathbf{r}^{\text{ref}}$. Let $(\hat{p}_{\tau b}, \hat{R}_{\tau b})$ denote a candidate’s realized link pose at horizon step τ , and let $(p_{\tau b}^{\text{ref}}, R_{\tau b}^{\text{ref}})$ be the corresponding desired pose. The objective is

$$J = \sum_{\tau=1}^H \omega_{\tau} \sum_{b \in \mathcal{B}} \eta_b \left[\|\hat{p}_{\tau b} - p_{\tau b}^{\text{ref}}\|_2^2 + 0.1 \text{ angle} \left(\hat{R}_{\tau b}^{\top} R_{\tau b}^{\text{ref}} \right)^2 \right], \quad (\text{A.6})$$

$$\omega_{\tau} = \frac{H - \tau + 1}{\sum_{j=1}^H (H - j + 1)}.$$

The body set $\mathcal{B}$ contains the torso and bilateral wrists. `angle` gives the relative rotation angle in radians. Link weights are $\eta_{\text{torso}} = 0.25$ and $\eta_{\text{left wrist}} = \eta_{\text{right wrist}} = 1$, without normalization across links. The pelvis has zero weight; joint, velocity, reference-deformation, and terminal penalties are inactive for this stage. The decreasing time weights prioritize the prefix that will be executed next. Scoring the realized motion against the unchanged target ensures that changing the input command alone cannot be counted as improved tracking.

Diagonal CEM and simulator failures. CEM (Pinneri et al., 2021) maintains a mean μ and a coordinate-wise standard deviation σ , initialized to zero and 1.6, respectively. Each iteration samples N latent vectors,

$$z^{(n)} = \mu + \sigma \odot \varepsilon^{(n)}, \quad \varepsilon^{(n)} \sim \mathcal{N}(0, I_d), \quad (\text{A.7})$$

then decodes and evaluates the corresponding commands in parallel rollouts. For the E lowest-cost candidates indexed by $\mathcal{E}$, the proposal update is

$$\mu_i^+ = \frac{1}{E} \sum_{n \in \mathcal{E}} z_i^{(n)},$$

$$\sigma_i^+ = \max \left\{ 0.075, \sqrt{\frac{1}{E} \sum_{n \in \mathcal{E}} \left(z_i^{(n)} - \mu_i^+ \right)^2} \right\}. \quad (\text{A.8})$$

The divisor is E , giving the population standard deviation of the elite coordinates. The proposal is diagonal, so the update does not estimate a full covariance from a small elite set. Sampling operates in normalized latent coordinates; physical clipping is applied by the decoder.

The rejection and validity terminology in Sec. 3.2 is implemented through failure-cost penalization for numerically invalid rollouts. Non-finite simulator measurements or objective values receive a large finite cost before CEM ranks the candidates. The elite update sorts these costs without a separate hard validity predicate; failed candidates are suppressed when sufficiently many lower-cost candidates are available, but are not unconditionally excluded. This mechanism does not certify general kinematic, contact, or physical feasibility.

Receding-horizon execution. One replanning update proceeds as follows:

1. Construct the next reference horizon and initialize the diagonal proposal. All candidates share the current simulated starting state.
2. Sample, decode, and roll out N sparse-command candidates; score them against the fixed target and update the proposal using the E lowest costs. Repeat for I updates.
3. Compare the nominal command, the best sampled candidate, and the final proposal mean using their rollout costs, and retain the lowest-cost choice.
4. Execute its first $\Delta = 10$ steps, record the realized configurations, and continue planning from the resulting simulator state.

At 50 Hz, the horizon covers 0.80 s and each retained prefix covers 0.20 s. The candidate count N denotes sampled proposals per iteration; nominal and final-mean evaluations are distinct comparison candidates. Concatenating executed prefixes yields $q_{1:T}^{\mathcal{R},\text{SCO}}$. Advancing from the realized state carries tracking deviations into the next optimization instead of restarting each horizon at an ideal reference pose.

A.3 Temporal smoothing and constrained arm refinement

Filtering on the rotation manifold. Savitzky–Golay filtering (Schaffer, 2011) smooths the three-point reference positions in the common Z-up world frame, using a five-frame window and a quadratic polynomial under the default settings. Orientation filtering uses local tangent coordinates so that smoothing does not produce invalid rotation matrices. For a window around frame t , the neighboring rotations R_j are represented relative to R_t as $\text{Log}(R_t^\top R_j)^\vee$, using the principal rotation vector with norm at most π . A local polynomial fit using the actual timestamps is evaluated at the timestamp of frame t to obtain $\delta\omega_t^{\text{SG}}$, and the filtered orientation is

$$R_t^{\text{SG}} = R_t \text{Exp}([\delta\omega_t^{\text{SG}}]_\times). \quad (\text{A.9})$$

The superscript $\vee$ converts a skew-symmetric matrix to its rotation vector. Near a segment boundary, the full window shifts inside the segment and the fitted polynomial is still evaluated at the target frame’s timestamp. The implementation splits the sequence where an adjacent timestamp gap exceeds three times the median sampling interval of the full sequence; segments shorter than five frames are left unchanged. Position filtering uses the same actual-time polynomial fit and boundary handling. Position and orientation filtering therefore use the same temporal reference while respecting their different geometry. Kimodo (Rempe et al., 2026) provides an initial whole-body configuration from the first sparse reference frame; the checked configuration initializes whole-body simulation.

Arm-only refinement and final assembly. SCO reduces sparse-reference error through the controller’s dynamics, but its realized wrists can still deviate from the retargeted manipulation reference. Temporal Mink IK (Zakka, 2024) addresses this residual by adjusting the arm chain, including the wrists. The optimized floating base, torso, and lower-body trajectory are held fixed, and the contact-aware finger trajectory is retained. Joint limits, per-step displacement and velocity bounds, and self-collision constraints restrict the arm updates. A posture penalty discourages unnecessary departures from the SCO configuration. Temporal restrictions are applied across successive refined states so that reducing a frame’s wrist error does not permit an unconstrained joint jump.

The final configuration combines the refined arms, retargeted fingers, and fixed non-arm motion. It supplies both the canonical motion reference and the body trajectory used for robot-view rendering. Holding the non-arm configuration fixed preserves those kinematic tracks; changes to the arms can still affect the response of a whole-body controller. The replay comparison in Sec. 4.5 therefore evaluates trackability before and after refinement. These measurements assess the demonstrated trajectories under the evaluated controller, without assuming that inverse kinematics alone establishes dynamic feasibility for arbitrary motions or deployment interfaces.

B Coordinate Conventions and Observation Synthesis

This appendix expands the geometric operations in Sec. 3.3. The recovered human scene and the rendered robot must be expressed in the same target-camera frame before their depths can be compared. The target

camera is obtained from the canonical whole-body motion, so scene reprojection and robot rendering use one pose at each time index.

B.1 Reference frames and the canonical record

Let $T_{A \leftarrow B} \in \text{SE}(3)$ map homogeneous coordinates in frame B into frame A . The world frame is denoted by W , the robot torso frame by B , and its camera frame by C . A camera-to-world pose has the form

$$C_t = \begin{bmatrix} R_t & p_t \\ 0^\top & 1 \end{bmatrix}, \quad C_t^{-1} = \begin{bmatrix} R_t^\top & -R_t^\top p_t \\ 0^\top & 1 \end{bmatrix}. \quad (\text{B.1})$$

Here $R_t \in \text{SO}(3)$ maps camera-frame vectors to the world frame, and $p_t \in \mathbb{R}^3$ is the camera origin expressed in world coordinates. Thus C_t maps camera-frame points to world coordinates, whereas C_t^{-1} maps world points into the camera frame for projection. All depths and translations in the geometric stages use a consistent metric scale. Image coordinates and calibrated intrinsics must refer to the same image size. The control clock provides the common time index for source observations, recovered geometry, canonical configurations, and the synthesized observations stored with them.

Write $T_{B \leftarrow C}$ for the fixed camera-to-torso transform from the target robot’s calibration and URDF. The sparse torso reference derived from human camera motion and the camera pose derived from canonical robot motion use this transform in opposite directions:

$$\begin{aligned} T_{W \leftarrow B}^{\text{ref}}(t) &= C_t^h T_{B \leftarrow C}^{-1}, \\ C_t^{\mathcal{R}} &= T_{W \leftarrow B}(q_t^{\mathcal{R}}) T_{B \leftarrow C}. \end{aligned} \quad (\text{B.2})$$

Here $T_{W \leftarrow B}(q_t^{\mathcal{R}})$ is the torso-to-world transform obtained by forward kinematics at the canonical robot configuration. The first line defines the torso target supplied to motion generation; the second uses the generated configuration to define the rendering camera. They need not yield identical camera paths because whole-body tracking and arm refinement construct the robot trajectory under its kinematic and control constraints. In particular, the observation stage uses the second line rather than substituting the source camera path or a constant vertical displacement.

The canonical record retains $(I_t^{\mathcal{R}}, q_t^{\mathcal{R}}, \ell)$ for a fixed target embodiment and camera convention. Its independence from the deployment interface means that these quantities are retained when the subsequent state–action mapping changes between SONIC and HEFT. A tracking controller is still used to construct the canonical motion. The body state recorded during deployment-controller replay can differ from the canonical configuration; the interface appendix specifies this distinction in detail.

B.2 Metric back-projection and target-camera visibility

Let K_h and $K_{\mathcal{R}}$ denote the source and target intrinsics. A source pixel is written as $u = (u_x, u_y)$, with homogeneous coordinate $\bar{u} = (u_x, u_y, 1)^\top$. EfficientSAM3 (Zeng et al., 2025) provides the source arm/hand mask. Invalidating the masked depth before projection prevents the source limbs from entering the reconstructed scene. For a remaining valid sample, optical-axis depth gives

$$\begin{aligned} p_h(u) &= \tilde{D}_t^h(u) K_h^{-1} \bar{u}, \\ \begin{bmatrix} p_{\mathcal{R}}(u) \\ 1 \end{bmatrix} &= (C_t^{\mathcal{R}})^{-1} C_t^h \begin{bmatrix} p_h(u) \\ 1 \end{bmatrix}, \\ v_x(u) &= f_x^{\mathcal{R}} \frac{[p_{\mathcal{R}}(u)]_x}{[p_{\mathcal{R}}(u)]_z} + c_x^{\mathcal{R}}, \quad v_y(u) = f_y^{\mathcal{R}} \frac{[p_{\mathcal{R}}(u)]_y}{[p_{\mathcal{R}}(u)]_z} + c_y^{\mathcal{R}}. \end{aligned} \quad (\text{B.3})$$

Here $\tilde{D}_t^h$ is the source depth after invalidating the human arm/hand mask. The points $p_h(u)$ and $p_{\mathcal{R}}(u)$ are expressed in the source-camera and target-camera frames, respectively; $v(u) = (v_x(u), v_y(u))$ is the projected target-image location. The quantities f_x, f_y are focal lengths in pixels and c_x, c_y are principal-point coordinates. The source depth is a z coordinate, rather than the Euclidean distance from the camera center. After transformation, the target depth used for visibility is $[p_{\mathcal{R}}(u)]_z$, rather than the original source depth.

Valid projected samples have finite coordinates and positive target z and contribute only inside the target image.

Forward splatting maps source colors into the target grid. Let $\mathcal{S}_t(v)$ denote the valid source samples whose projected splats cover target pixel v . The metric z-buffer selects the nearest target-camera surface,

$$D_t^{\text{geo}}(v) = \min_{u \in \mathcal{S}_t(v)} [p_{\mathcal{R}}(u)]_z, \quad V_t^{\text{geo}}(v) = \mathbf{1}[\mathcal{S}_t(v) \neq \emptyset]. \quad (\text{B.4})$$

The minimum is used only where $V_t^{\text{geo}} = 1$; elsewhere the depth is invalid. Color is taken from a sample at the nearest surface, with numerical ties resolved consistently. This visibility ordering prevents a farther surface from overwriting a nearer one merely because it is processed later. The geometric hole mask is $H_t^{\text{geo}} = 1 - V_t^{\text{geo}}$. It includes areas exposed by source-arm removal and areas that were not observed from the original viewpoint. Reprojection changes where visible source evidence is placed; it does not supply the appearance of these newly exposed regions.

B.3 Robot rendering, compositing, and restricted completion

MuJoCo renders robot color and depth from $q_t^{\mathcal{R}}$ with the same $C_t^{\mathcal{R}}$ used in Eq. B.3. The rendered depth is represented in the same metric target-camera convention before visibility comparison. Write V_t^{rob} for its binary validity mask. The visible-robot mask is defined at every target pixel by

$$M_t^{\text{rob}}(v) = \begin{cases} \mathbf{1}[D_t^{\text{rob}}(v) < D_t^{\text{geo}}(v)], & (V_t^{\text{rob}}(v), V_t^{\text{geo}}(v)) = (1, 1), \\ 1, & (V_t^{\text{rob}}(v), V_t^{\text{geo}}(v)) = (1, 0), \\ 0, & (V_t^{\text{rob}}(v), V_t^{\text{geo}}(v)) = (0, 1), \\ 0, & (V_t^{\text{rob}}(v), V_t^{\text{geo}}(v)) = (0, 0). \end{cases} \quad (\text{B.5})$$

If only one surface is valid, that surface supplies the pixel. When neither is available, the pixel remains a hole. Neighboring valid scene depths may supply local depth context inside geometric holes, as described in the main pipeline; they do not replace the depth comparison at pixels with measured scene support. The visible-robot mask, rather than the rendered silhouette alone, controls the composite:

$$I_t^{\text{comp}} = M_t^{\text{rob}} I_t^{\text{rob}} + (1 - M_t^{\text{rob}}) I_t^{\text{geo}}, \quad H_t^{\text{rem}} = H_t^{\text{geo}} (1 - M_t^{\text{rob}}). \quad (\text{B.6})$$

Products here are pixel-wise, with binary masks broadcast over color channels. The color arrays I_t^{rob} and I_t^{geo} use finite zero placeholders wherever their respective validity masks are zero, so masked arithmetic is defined even at holes. A branch-based selection of the valid color gives the same composite. In particular, $H_t^{\text{rem}} = (1 - V_t^{\text{geo}})(1 - V_t^{\text{rob}})$: neither source supplies a pixel precisely when both validity masks are zero. A foreground scene object can occlude the robot, and a nearer robot surface can occlude the reprojected scene. The second expression is the mask form of the residual-hole set difference in the main text: only pixels uncovered by both sources remain to be completed.

DiffuEraser (Li et al., 2025) supplies appearance in H_t^{rem} . Denoting its completed prediction by I_t^{fill} , the final observation can be expressed as

$$I_t^{\mathcal{R}} = (1 - H_t^{\text{rem}}) I_t^{\text{comp}} + H_t^{\text{rem}} I_t^{\text{fill}}. \quad (\text{B.7})$$

Thus, valid scene and rendered-robot pixels remain fixed. Completion has a restricted appearance role after geometric visibility has been established; it does not change the canonical configuration, target camera pose, or action labels. The completed observation is stored with the same $q_t^{\mathcal{R}}$ and instruction to form the canonical episode.

For the reported 400 egocentric training clips, source depth is recorded by the iPhone. The RGB-only route in the framework instead uses metric depth from Depth Anything 3 (Lin et al., 2025) before the same geometric operations. These are alternative sources for D_t^h , rather than two separately evaluated training-data conditions. The observation ablation in Sec. 4.3 compares how views are aligned while keeping the converted motion and state-action supervision fixed.

Table C.1 Representations used in control-command generation. Dimensions refer to one knot or one time-indexed record, as indicated. The search coordinates are shared, while body-action labels depend on the controller.

Quantity	Dimension	Contents and role
Search residual	35 per knot	Pelvis translation and rotation-vector corrections, plus 29 body-joint corrections.
Horizon search	350	Ten residual knots interpolated over the horizon.
SONIC body label	64	Motion token from the optimized motion reference.
HEFT body label	$38 = 29 + 9$	Absolute body-joint targets and pelvis code.
Hand action	12	Bilateral BrainCo targets, with grasp and release handled separately.
Complete action	$76 / 50$	SONIC / HEFT body label with 12 hand targets.
Policy state	$44 = 29 + 3 + 12$	Rollout body joints, base-frame gravity, and retargeted hand angles.

C Controller-Specific Command and State Interfaces

This appendix expands Sec. 3.4. For each deployment controller, KODA optimizes a body reference against the fixed canonical trajectory, evaluates the realized motion, and records state-action supervision. Table C.1 distinguishes the search representations from the quantities exported for VLA training.

C.1 Shared Body-Reference Search

Let m_t^{can} contain the canonical pelvis pose, the 29-dimensional body-joint vector b_t^{can} , and the associated link poses. The reference is represented on the 50 Hz control clock; joint velocities and link quantities are available with the motion record. We use $\tilde{m}_t$ for a candidate command reference and $\hat{m}_t^\kappa$ for the motion realized by controller κ . These three quantities serve different roles: the canonical motion specifies the desired behavior, the candidate is supplied through the controller’s reference interface, and the rollout determines its tracking cost.

For knot index $k = 1, \dots, K$, each temporal knot contains normalized residuals

$$z_k = (z_k^p, z_k^\omega, z_k^b) \in \mathbb{R}^{3+3+29}, \quad z = (z_1, \dots, z_K), \quad K = 10. \quad (\text{C.1})$$

Here $z_k^p, z_k^\omega \in \mathbb{R}^3$ and $z_k^b \in \mathbb{R}^{29}$ are dimensionless pelvis-translation, rotation-vector, and body-joint corrections, respectively. The resulting search has 350 coordinates for either controller. Hands are excluded from this body-reference search and retain the retargeted geometric trajectory used for the hand-state and hand-command construction below. Let $S_\kappa \in \mathbb{R}^{35 \times 35}$ denote the diagonal conversion from normalized coordinates to physical translation, rotation, and joint corrections. Both controllers use the same scales:

$$S_{\text{SONIC}} = S_{\text{HEFT}} = \text{diag}(0.055I_2, 0.0275, 0.0825I_3, 0.1375I_{29}),$$

where I_d is the d -dimensional identity matrix. The first three scale entries are in meters; the rotation and joint entries are in radians. For control-step index τ within the current horizon, let $\delta_\tau = (\delta p_\tau, \delta \omega_\tau, \delta b_\tau) \in \mathbb{R}^{35}$ collect the dense physical corrections. With interpolation weights $A_{\tau k}$, they are

$$\delta_\tau = \sum_{k=1}^K A_{\tau k} S_\kappa \text{clip}(z_k, -1, 1), \quad \sum_k A_{\tau k} = 1. \quad (\text{C.2})$$

The knots are uniformly spaced over the horizon, and the nonnegative weights $A_{\tau k}$ linearly interpolate neighboring knots. The clipping operation acts elementwise. Let p_τ^{can} and R_τ^{can} be the canonical pelvis position and pelvis-to-world rotation. Translation and rotation-vector corrections are expressed in the world frame and applied as

$$\begin{aligned} \tilde{p}_\tau &= p_\tau^{\text{can}} + \delta p_\tau, \\ \tilde{R}_\tau &= \text{Exp}([\delta \omega_\tau]_\times) R_\tau^{\text{can}}, \\ \tilde{b}_\tau &= b_\tau^{\text{can}} + \delta b_\tau, \end{aligned}$$

where $[v]_{\times} w = v \times w$ and Exp maps a skew-symmetric matrix to a rotation. Body-joint corrections are followed by elementwise clipping to joint limits when those limits are configured. Forward kinematics then recomputes the link poses from the perturbed pelvis and body joints, maintaining a single connected robot configuration.

The controller adapter encodes each candidate for simulation using SONIC motion encoding or the HEFT representation below. CEM searches the common physical reference space and evaluates each candidate through the selected controller. Its 350 search coordinates differ from the 64-dimensional SONIC token and 38-dimensional HEFT body label.

C.2 Tracking Cost and Receding-Horizon Settings

The shared objective evaluates 29 body joints and 14 links: the pelvis, torso, and bilateral hip, knee, ankle, shoulder, elbow, and wrist links. For two motion states m, m' , let $b, b' \in \mathbb{R}^{29}$ be their body-joint angle vectors. Let $p_j, p'_j \in \mathbb{R}^3$ be the world-frame positions of link j , and $R_j, R'_j \in SO(3)$ its link-to-world rotations, for $j = 1, \dots, 14$. Define

$$\begin{aligned} e_q(m, m') &= \frac{1}{29} \|b - b'\|_2^2, \\ e_p(m, m') &= \frac{1}{14} \sum_{j=1}^{14} \|p_j - p'_j\|_2^2, \\ e_R(m, m') &= \frac{1}{14} \sum_{j=1}^{14} \theta(R_j^\top R'_j)^2, \\ \theta(Q) &= \arccos\left(\text{clip}\left(\frac{\text{tr}(Q) - 1}{2}, -1, 1\right)\right). \end{aligned} \tag{C.3}$$

Here $Q \in SO(3)$ denotes a relative rotation and $\theta(Q) \in [0, \pi]$ its geodesic angle. Positions are in meters and angular quantities are in radians. Each link contributes equally to the link averages; joint error is averaged over the 29 joint coordinates. The orientation term is the squared geodesic angle of the relative rotation. Its quaternion implementation uses the absolute quaternion inner product, so the two quaternion signs representing the same rotation produce the same error.

For a horizon of H steps, the complete objective is

$$\begin{aligned} J_\kappa &= \sum_{\tau=1}^H \left[e_q(\hat{m}_\tau^\kappa, m_\tau^{\text{can}}) + e_p(\hat{m}_\tau^\kappa, m_\tau^{\text{can}}) + 0.5e_R(\hat{m}_\tau^\kappa, m_\tau^{\text{can}}) \right] \\ &\quad + \mathcal{L}_{\text{def}} + 0.1 \sum_{\tau=2}^H \|u_\tau^{\text{PD}} - u_{\tau-1}^{\text{PD}}\|_2^2 + \mathcal{L}_{\text{terminal}}, \\ \mathcal{L}_{\text{def}} &= \sum_{\tau=1}^H \left[0.1e_q(\tilde{m}_\tau, m_\tau^{\text{can}}) + 0.1e_p(\tilde{m}_\tau, m_\tau^{\text{can}}) + 0.05e_R(\tilde{m}_\tau, m_\tau^{\text{can}}) \right], \\ \mathcal{L}_{\text{terminal}} &= 5e_q(\hat{m}_H^\kappa, m_H^{\text{can}}) + 5e_p(\hat{m}_H^\kappa, m_H^{\text{can}}). \end{aligned} \tag{C.4}$$

The tracking terms compare realized motion with the fixed canonical target. The deformation terms instead compare the candidate reference with that target, penalizing unnecessary changes to the commands. The action-rate term uses $u_\tau^{\text{PD}} \in \mathbb{R}^{29}$, the low-level PD position targets produced during rollout for the 29 body joints, in radians. Hand targets are excluded. This term sums squared changes across joints and has zero contribution at the first horizon step. It does not measure changes between motion tokens or between the normalized search vectors.

All running terms use uniform time weights of one. The horizon reduction is a sum, while the joint and link reductions in Eq. C.3 are averages. The terminal penalties add weight to joint and link-position errors at the last step; no terminal orientation term is included in this objective.

Control Command Generation uses $N = 80$ candidates per CEM iteration, $I = 8$ iterations per replan, and $E = 9$ elites. The planning horizon is $H = 40$ steps, or 0.80 s, and the retained prefix is $\Delta = 20$ steps, or

0.40s. After executing that prefix in simulation, the next search continues from its resulting state against the subsequent canonical reference window. Candidate rollouts with non-finite joint, link, or action values receive the failure cost 10^{12} . The objective otherwise scores the realized motion through the stated tracking terms; it does not add a separate fall or contact heuristic. These settings apply to the second optimization stage for both evaluated deployment controllers.

At the end of a sequence, the reference window is padded by repeating its last frame to retain the full planning horizon. If n_{rem} control steps remain, only $\min(\Delta, n_{\text{rem}})$ steps are executed.

C.3 SONIC and HEFT Body-Action Labels

SONIC’s motion encoder maps the optimized reference to a 64-dimensional motion token (Luo et al., 2026). This token is recorded as the body-action label associated with the selected reference, whereas rollout body-joint angles describe the resulting motion.

HEFT uses 29 absolute body-joint targets together with a nine-dimensional pelvis code (Liu et al., 2026a). Let $p_t \in \mathbb{R}^3$ and $(\alpha_t, \beta_t, \psi_t)$ denote the optimized pelvis position and roll, pitch, and yaw, with position expressed in the world frame. Let $R_z(\psi)$ be the rotation about the world z axis that maps vectors from the yaw-aligned frame into the world frame. Define

$$\begin{aligned} c_t &= (\sin \alpha_t, \cos \alpha_t - 1, \sin \beta_t, \cos \beta_t - 1, \delta\psi_t, \delta p_t^\top, p_{t,z})^\top, \\ \delta\psi_t &= \text{wrap}(\psi_{t+1} - \psi_t), \quad \delta p_t = R_z(\psi_t)^\top (p_{t+1} - p_t). \end{aligned} \tag{C.5}$$

The row combines the current pose’s roll, pitch, and absolute height with the transition to the next pose. The translation increment is expressed in the current yaw frame and contains three coordinates. A valid row therefore needs a successor pelvis pose; this is part of the motion-code definition rather than an extra body degree of freedom. Encoding $L + 1$ consecutive pelvis poses produces L code rows; the final input pose supplies the successor needed by the last row.

During decoding, planar position and heading accumulate from an anchor $(\bar{p}_{0,xy}, \bar{\psi}_0)$. Let L be the number of code rows and use local row indices $i = 0, \dots, L - 1$, independent of the episode time index t . With code components numbered from one, set $\delta\psi_i = c_{i,5}$ and $\delta p_i = (c_{i,6}, c_{i,7}, c_{i,8})^\top$. Then

$$\begin{aligned} \bar{\psi}_i &= \bar{\psi}_0 + \sum_{j=0}^{i-1} \delta\psi_j, \\ \bar{p}_{i,xy} &= \bar{p}_{0,xy} + \sum_{j=0}^{i-1} [R_z(\bar{\psi}_j) \delta p_j]_{xy}, \quad \bar{p}_{i,z} = c_{i,9}, \\ \bar{\alpha}_i &= \text{atan2}(c_{i,1}, 1 + c_{i,2}), \quad \bar{\beta}_i = \text{atan2}(c_{i,3}, 1 + c_{i,4}). \end{aligned} \tag{C.6}$$

Empty sums are zero, and code components are numbered from one. Absolute height directly defines the decoded vertical position; it is not obtained by cumulatively adding the vertical increment. The final row’s transition leads beyond the returned sequence of current poses. This explains how the code can contain both vertical displacement and absolute height without integrating both into the same position.

C.4 Hand Targets and Recorded State

Both body-action formats append 12 BrainCo hand-joint targets, six per hand, yielding 76-dimensional SONIC actions and 50-dimensional HEFT actions. The hand command is constructed from the retargeted robot hand geometry. Following EgoHumanoid (Shi et al., 2026), the mean finger curvature is thresholded separately for each hand to distinguish grasp from release. For grasp, the command uses the calibrated BrainCo grasp pose to encourage active closure. For release, it uses the current retargeted joint angles. These targets are delivered separately to hand control. This construction specifies the closure command without interpreting a retargeted geometric pose as a measured grasp force.

The policy state has the same 44-dimensional format for both controllers:

$$s_t^\kappa = \left(b_t^{\text{roll}, \kappa}, (R_t^{\text{base}, \kappa})^\top g^W, f_t^{\text{ret}} \right). \tag{C.7}$$

Here $b_t^{\text{roll}, \kappa} \in \mathbb{R}^{29}$ contains rollout body-joint angles and $f_t^{\text{ret}} \in \mathbb{R}^{12}$ contains continuous retargeted hand angles, both in radians. The rotation $R_t^{\text{base}, \kappa} \in SO(3)$ maps base-frame vectors to the world frame, and $g^W = (0, 0, -1)^\top$ is the dimensionless unit gravity direction. Thus $(R_t^{\text{base}, \kappa})^\top g^W$ gives the three-dimensional projected-gravity observation in the rollout base frame. The state records the estimated hand geometry even when a grasp action requests the calibrated closure target. Replacing the state hand angles with that target would change the meaning of the observation. Body-state values depend on the selected controller’s rollout, although the state format and retained hand estimates are shared.

C.5 Episode Export and Observation Reuse

The procedure below summarizes the conversion from a canonical episode to controller-specific supervision.

Controller-specific episode export

1. Read the canonical motion, synthesized images, retargeted hand trajectory, and instruction on their common temporal index.
2. For each receding-horizon window, decode CEM residuals into candidate body references, encode them for controller κ , simulate, and evaluate Eq. C.4 against the fixed canonical target.
3. Retain the selected prefix of up to 20 steps, shortened when fewer steps remain, and continue from its resulting simulated state. Collect the optimized reference and realized body motion.
4. Record the selected reference as SONIC tokens or HEFT body labels; append the 12 grasp/release hand targets.
5. Assemble the 44-dimensional state from rollout body joints, projected gravity, and retargeted hand angles.
6. Associate these states and actions with the retained canonical images and instruction to obtain the controller-specific training episode.

Canonical-image reuse follows the fixed reference trajectory, while the recorded body state follows the selected controller’s realized motion. Their correspondence is therefore approximate when tracking is imperfect. The wrist and camera deviations in TABLE 4 measure this residual mismatch for the evaluated controllers, and TABLE 1 evaluates the utility of the resulting episodes for co-training. Changing from SONIC to HEFT repeats command optimization, encoding, and state recording while retaining canonical motion generation and view synthesis. The reuse rests on tracking the shared reference sufficiently well for the evaluated setting; it does not assert identical camera poses across controller rollouts.

D Experimental Protocol and Metric Definitions

This appendix specifies the data composition, controlled comparisons, and interpretation of the evaluation metrics used in Sec. 4.1. The task-success results remain in TABLES 1–3; the motion and rollout errors remain in TABLE 4.

D.1 Demonstration Sources and Evaluation Conditions

The experimental platform is a Unitree G1 with BrainCo dexterous hands. For each task and deployment controller, the robot dataset contains 50 demonstrations collected through the corresponding SONIC or HEFT PICO teleoperation pipeline. The human pool contains 100 demonstrations per task, recorded with a chest-mounted iPhone using MobileEgo Anywhere (Palanisamy et al., 2026). All 400 human clips include phone-recorded depth, and both demonstration sources use 50 Hz sequences. The same human canonical records supply the motion, synthesized observations, and instructions for both controllers; command generation and state recording produce the controller-specific training fields.

TABLE D.1 specifies the task goals and condition changes. Robot-ID is defined relative to object appearance and pose conditions in the robot demonstrations. Robot-OOD changes these conditions beyond that robot dataset. Each task’s human pool contains 80 shifted and 20 robot-matched clips, so the shifted test conditions are represented in the human training pool. This evaluates whether converted human data supplies coverage

Table D.1 Task goals and the changes defining Robot-OOD. Robot-ID uses the appearance and pose conditions represented in the robot demonstrations.

Task	Placement goal	Robot-ID to Robot-OOD change
Cup-to-basket	Place the object in a basket	Green bottle replaced by a brown bottle
Parcel-to-shelf	Place the parcel on a shelf	Purple parcel replaced by a white parcel
Basket-to-shelf	Place the basket on a shelf	Basket handle tilted by approximately 30°
Cup-to-drawer	Place the object in a drawer	Green bottle replaced by a brown bottle

absent from the robot demonstrations. It does not test conditions excluded from both training sources. The 80:20 composition describes the complete human pool; the independently sampled smaller budgets are not specified as having that same proportion.

D.2 Training Budgets and Success-Rate Aggregation

For each controller configuration, GR00T N1.7 (Bjorck et al., 2025) is trained jointly on all four tasks for 20,000 iterations with batch size 160. Within each task, sampling from robot and human sources is proportional to episode counts; other sampling settings follow GR00T defaults. The scaling study uses $k \in \{0, 25, 50, 75, 100\}$ converted human demonstrations per task, with the 50 robot demonstrations held fixed. The human subset is randomly sampled independently at each budget without nesting. Thus, intermediate budgets differ in both the amount and identity of the available human data.

The source proportions implied by the episode counts are

$$p_{\text{robot}}(k) = \frac{50}{50 + k}, \quad p_{\text{human}}(k) = \frac{k}{50 + k}. \quad (\text{D.1})$$

At the full budget, the robot-to-human episode ratio is 1 : 2. These proportions describe source sampling within a task; they do not specify the realized source counts of each individual minibatch. Increasing k changes the training data mixture while preserving the policy architecture and iteration budget. The component ablations all use SONIC with 50 robot and 100 converted-human demonstrations per task.

Each configuration is trained once and evaluated in 20 real-robot trials per task and condition. A trial succeeds when the placement is completed without a fall or unrecovered manipulation failure. Let $y_{j,c,i} \in \{0, 1\}$ denote success in trial i for task j and condition $c \in \{\text{Robot-ID}, \text{Robot-OOD}\}$. The reported task success rate and macro-average, in percent, are

$$S_{j,c} = \frac{100}{20} \sum_{i=1}^{20} y_{j,c,i}, \quad \bar{S}_c = \frac{1}{4} \sum_{j=1}^4 S_{j,c}. \quad (\text{D.2})$$

The four tasks receive equal weight. A single additional success changes a task rate by 5 percentage points and its corresponding macro-average by 1.25 points. Because each task has the same trial count, the macro-average also equals the proportion of successes among the 80 trials for that condition. This arithmetic resolution does not quantify variation across training runs: the results are point estimates from one trained policy per configuration, rather than averages over repeated training seeds.

D.3 Controlled Component Comparisons

TABLE D.2 separates the observation variants from the control-command variants. In the observation study, the demonstration set, state-action supervision, policy architecture, training budget, and real-robot evaluation protocol are fixed. Only the human-derived observation construction changes. Fixed offset adapts the view-alignment strategy of EgoHumanoid (Shi et al., 2026) to chest-mounted capture, using a 0.10 m vertical shift and ± 0.05 m training-time perturbations. Both reprojection-only variants include hole inpainting.

Table D.2 Definitions of the component variants. All comparisons use SONIC and the full human-data budget. The table describes changed components; quantitative policy results are in TABLES 2 and 3.

Variant	Observation or command construction
<i>Observation alignment: state-action supervision held fixed</i>	
No alignment	Original human images.
Fixed offset	Scene reprojection with a fixed vertical camera shift, training-time perturbations, and hole inpainting.
Motion-based reprojection	Scene reprojection and hole inpainting using camera poses derived from the canonical robot motion.
Full alignment	Motion-derived onboard views, additionally removing human arms and rendering and compositing the target robot.
<i>Control Command Generation: canonical trajectories and images held fixed</i>	
Without SCO	Canonical whole-body motion used directly as the tracking controller’s reference; command optimization omitted only at this stage.
With SCO	Controller-specific commands optimized to track the same canonical motion.

Full alignment changes human-arm removal and robot rendering/compositing together relative to motion-based reprojection. Consequently, this comparison does not separately identify the effect of each image-processing operation. The main interpretation concerns the complete alignment pipeline relative to original human observations and the adapted fixed-offset baseline. The intermediate variant describes a partial pipeline under the same training setup.

In the command-generation study, both variants retain Whole-Body Motion Generation and Robot Onboard View Generation. They use identical canonical trajectories, observations, demonstrations, and policy-training settings. The optimized commands are evaluated through the deployment controller, and body-state features are recorded from the resulting rollouts. Changing command generation can therefore change both action labels and body-state features. This is a comparison of state-action supervision as a whole, not an action-only intervention with all state inputs held fixed. These component comparisons use SONIC; the data-scaling study evaluates reuse of the complete converted records under both SONIC and HEFT.

D.4 Motion References, Units, and Refinement Replay

The blocks of TABLE 4 use different references. Whole-body motion-generation errors measure agreement with the source motion references. Controller-rollout errors measure agreement between realized motion and the fixed canonical trajectory. The former characterizes source-reference fidelity; the latter characterizes realization by a particular controller. Direct denotes unoptimized references, SCO denotes CEM-optimized commands, and SCO+IK adds Mink arm refinement during whole-body motion generation.

Position errors use centimeters and orientation errors use radians. The motion-generation block reports wrist position and orientation and torso orientation; it does not report a torso-position metric. The rollout blocks report wrist and camera position and orientation. Their minima are selected separately for each error type. In particular, improved position tracking does not imply improved orientation tracking: SONIC’s camera-orientation error is 0.0784 rad for Direct and 0.0808 rad for SCO. Dashes denote unreported metrics or variants, not zero error.

Motion-generation entries average aggregate errors from LAFAN (Harvey et al., 2020) and the 400 egocentric training clips. LAFAN provides an auxiliary check on highly dynamic motion, while the egocentric clips connect the diagnostics to the converted policy-training data. Both controllers share evaluation data. The aggregate entries are not separate results for the two sources, and their source-reference and rollout-reference errors should not be interchanged.

TABLE D.3 expands the before/after-Mink replay diagnostics described in Sec. 4.5. The same SONIC tracker is used to replay the pre-refinement and refined trajectories on both data sources. These diagnostics include pelvis-relative position errors, torso orientation, and non-arm joint-angle RMSE. They are distinct from the wrist errors measured against the original source motion. The small changes characterize comparable controller trackability on the evaluated trajectories; they do not constitute a separate downstream policy-success ablation of Mink refinement.

Table D.3 SONIC replay before and after Mink arm refinement on LAFAN and the 400 egocentric clips. These are the existing replay diagnostics, shown separately by metric and unit.

Replay metric	Unit	Before refinement	After refinement
Pelvis-relative non-arm position error	cm	2.60	2.73
Pelvis-relative wrist-position error	cm	1.59	1.75
Torso-orientation error	rad	0.0588	0.0579
Non-arm joint RMSE	rad	0.0750	0.0784

E Qualitative Demonstration Timelines

Each timeline spans one complete source clip through eight selected keyframes, including its first and last frames. Read left to right, then top to bottom; timestamps refer to the original 50 Hz overlay video. Each frame retains the human view (upper left), synthesized robot view (upper right), and offline MuJoCo kinematic replay (main panel). These illustrate converted demonstrations; they are not additional real-robot evaluation trials.




Figure E.1 Complete-clip timeline: bimanual box transfer. Selected moments from a 13.28s clip show reaching toward a box, lifting and carrying it, and setting it on a table before withdrawing the hands. Frames 1–8 preserve chronological order and the paired human, robot-view, and whole-body-motion displays.




Figure E.2 Complete-clip timeline: sequential object manipulation. Eight selected moments span a 30.76 s clip: a cup is moved toward a basket, a parcel is transferred to a shelf, and a bottle is carried to a refrigerator and placed inside before the door is closed. The first and last video frames are included. Time advances left to right and then top to bottom; the three synchronized displays have the same meanings as in Fig. E.1.

F Additional Task Demonstrations and Deployments

Figures F.1–F.3 complement Fig. 4 with the other three evaluated tasks. Each figure follows the same row order: Robot Data in the Robot-ID condition, converted KODA Data, and successful SONIC deployment in the Robot-OOD condition. Four keyframes show the action progression within each row. Timestamps are relative to the corresponding source clip; different rows are independent recordings and are not temporally synchronized. The KODA row retains the human view, synthesized robot view, and offline kinematic replay described in Appendix E. Bystanders are cropped out or pixelated.




Figure F.1 Cup-to-basket. Robot Data shows the green bottle used in Robot-ID. KODA Data covers the brown-bottle condition, and the deployment row shows a successful Robot-OOD placement into the basket. Each row progresses from grasping through transfer to placement.




Figure F.2 Parcel-to-shelf. The purple parcel in Robot Data is replaced by a white parcel in the KODA and Robot-OOD deployment examples. The keyframes show reaching, lifting and transferring the parcel, followed by placement on the shelf.



Figure F.3 Basket-to-shelf. The Robot-OOD change is an approximately 30° tilt of the basket handle relative to the Robot-ID condition. KODA Data supplies a corresponding converted demonstration; the deployment row shows grasping the handle, lifting the basket, and placing it on the shelf.